

Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL:
 - SFW and its semantics
 - LIKE, AS, *, %, ...
 - Null values
 - Single and multiple relations
 - Subqueries
 - Bag algebra and set operations
 - **Next: grouping**



Aggregations

- SUM, MAX, MIN, COUNT, AVG can be applied to a column in a SELECT clause to produce that aggregation
- COUNT (*) is a special syntax to count number of tuples
- Number of films:
 - SELECT COUNT(*) FROM Film
- Max and average film lengths:
 - SELECT MAX(length), AVG (length) FROM Film



Removing Duplicates

- Adding DISTINCT removes duplicates prior to aggregation
 - `SELECT COUNT(DISTINCT length) FROM Film`
 - `SELECT COUNT(DISTINCT release_year) FROM Film`



NULL values

- Note:
 - NULL values are not involved in aggregations
 - But if there are no non-NULL values, result will be a NULL
- Number of movies
 - `SELECT COUNT (*) FROM Film;`
- Number of movies with a non-null length
 - `SELECT COUNT (length) FROM Film;`



Grouping

- In some cases, we may want to compute an aggregate for each “group” of tuples
- You do so by adding a GROUP BY clause after SELECT FROM WHERE
 - The results of SFV are then grouped according to the grouping attributes, and aggregation is applied per group
 - SELECT rating, AVG(length), MIN (length)
 - FROM Film GROUP BY rating;



Grouping (contd.)

- Film_actor (actor_id, film_id, last_update), Actor (actor_id, first_name, last_name, last_update), Film (film_id, ...)
 - SELECT Actor.first_name, Actor.last_name, COUNT(*)
 - FROM Film, Actor, Film_actor
 - WHERE Film.film_id = Film_actor.film_id
 - AND Actor.actor_id = Film_actor.actor_id
 - GROUP BY Actor.first_name, Actor.last_name



Restriction of SELECT list with aggregation

- If aggregation is used, then each element of the SELECT clause must either be:
 - An aggregate, or
 - An attribute in the GROUP-BY list
- Why this restriction?



Restriction of SELECT list with aggregation

- If aggregation is used, then each element of the SELECT clause must either be:
 - An aggregate, or
 - An attribute in the GROUP-BY list
- Why this restriction?
 - If an attribute is not being aggregated or being grouped, then you need some way to “squish” the values down per group.



Exercise

- Find the movie with the maximum length



Exercise

- Find the movie with the maximum length
 - `SELECT FI.title, FI.length`
 - `FROM Film AS FI`
 - `WHERE FI.length >= (SELECT MAX (length) FROM Film)`



HAVING Clauses

- HAVING <condition> may follow a GROUP BY clause
- If so, the condition is applied to each group, and groups not satisfying the condition are eliminated
- Example:
- Show actors who have starred in at least 30 movies
 - SELECT Actor.first_name, Actor.last_name, COUNT(*)
 - FROM Film, Actor, Film_actor
 - WHERE Film.film_id = Film_actor.film_id
 - AND Actor.actor_id = Film_actor.actor_id
 - GROUP BY Actor.first_name, Actor.last_name
 - **HAVING COUNT (*) > 30**



Restrictions on HAVING Clauses

- Similar to SELECT clauses: each attribute mentioned must either be part of the GROUP BY or be aggregated



General Form of GROUPING

```
SELECT S  
FROM R1, R2, ...  
WHERE C1  
GROUP BY A1, A2, ...  
HAVING C2
```

- S and C2 can contain A1, A2, ... or any other aggregated attributes
- C1: any condition



General Form of GROUPING

```
SELECT S  
FROM R1, R2, ...  
WHERE C1  
GROUP BY A1, A2, ...  
HAVING C2
```

- Order of evaluation:
 - Compute the “FROM-WHERE” part:
 - For each combination of tuples in the cross product of R1, R2, ...
 - Keep only those tuples that satisfy C1
 - Group by A1, A2, ...
 - For each group, check if C2 is satisfied
 - If so: compute aggregates in S and add to output



Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL:
 - SFW and its semantics
 - LIKE, AS, *, %, ...
 - Null values
 - Single and multiple relations
 - Subqueries
 - Bag algebra and set operations
 - Grouping
 - Next: Ordering



Ordering: ORDER BY

- Use an ORDER BY clause to enforce ordering of the result
- ORDER BY <attr> ASC | DESC
 - Order films by descending length:
 - SELECT title, length
 - FROM Film
 - ORDER BY length DESC;
 - Order films by descending length, then title ascending
 - SELECT title, length
 - FROM Film
 - ORDER BY length DESC, title ASC;



Restrict output: LIMIT & OFFSET

- Sometimes you want to limit the result to a few tuples
- LIMIT k
 - Order films by descending length, ascending title, return top 15
 - SELECT title, length FROM Film ORDER BY length DESC, title ASC
 - LIMIT 15
- And sometimes you want to start the output at a particular point
- OFFSET k
 - Order films by descending length, ascending title, return from position 11 to 15
 - SELECT title, length FROM Film ORDER BY length DESC, title ASC
 - LIMIT 5 OFFSET 10



Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL:
 - SFW and its semantics
 - LIKE, AS, *, %, ...
 - Null values
 - Single and multiple relations
 - Subqueries
 - Bag algebra and set operations
 - Grouping
 - Ordering
 - Case



Case Statements

- CASE statements allow you to create new derived data in the select clause
- SELECT title, CASE
 - WHEN length > 180 THEN 'long'
 - WHEN length > 120 AND length <= 180 THEN 'medium'
 - WHEN length < 120 AND length >= 60 THEN 'short'
 - ELSE 'super short'
 - END AS type
- FROM Film

