

So far: Querying with SQL

- Lots more functionality within SQL to support even more advanced querying
 - You can call functions in programming languages all within a **SELECT** statement
 - You can use “**WINDOW**” functions that allow you to perform computation across groups, e.g., capture a running/cumulative total
 - You can define special shortcuts for relations via Common Table Expressions (**CTEs**)
- We won't cover any of this for now...



Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL Queries:
 - SFV and its semantics
 - LIKE, AS, *, %, ...
 - Null values
 - Single and multiple relations
 - Subqueries
 - Bag algebra and set operations
 - Grouping
 - Ordering
 - Case
- Next: SQL Modifications



It's not sufficient to query...

- We also need to be able to **modify** the data in relations
- A modification command doesn't return the result in the same way a query does, but it changes the contents of the database
- Three kinds of modifications:
 - *Insert* a tuple or tuples
 - *Delete* an existing tuple (or tuples)
 - *Update* the value(s) of an existing tuple (or tuples)



Inserting One or More Tuples

- `INSERT INTO Relation VALUES (<list of values>);`
- Example:
 - Actor (actor_id, first_name, last_name, last_update)
 - `INSERT INTO Actor VALUES (5000, 'Adam', 'Driver', 'Feb 3, 2020');`
 - `INSERT INTO Actor VALUES (5001, 'Brenda', 'Driver', 'Feb 3, 2020'), (5002, 'Charles', 'Driver', 'Feb 3, 2020');`



Specifying Attributes in Insert

- `INSERT INTO Actor VALUES (5001, 'Brenda', 'Driver', 'Feb 3, 2020'), (5002, 'Charles', 'Driver', 'Feb 3, 2020');`
- The previous syntax requires you to remember order of attributes
- Can be more explicit:
 - `INSERT INTO Actor (actor_id, first_name, last_name, last_update) VALUES (5003, 'Didi', 'Driver', 'Feb 3, 2020');`
- Another benefit beyond remembering the order of attributes:
 - We don't have values for all attributes, and we want the system to fill in NULL values for the rest



Inserting Many Tuples

- We may insert the entire result of a query into a relation
- `INSERT INTO Relation (<Subquery>);`
- Actor (actor_id, first_name, last_name, last_update)
- `INSERT INTO Actor (actor_id) (SELECT * FROM Film_Actor)`
- Q: What do you think this query does?



Inserting Many Tuples

- We may insert the entire result of a query into a relation
- `INSERT INTO Relation (<Subquery>);`
- Actor (actor_id, first_name, last_name, last_update)
- `INSERT INTO Actor (actor_id) (SELECT * FROM Film_Actor)`
- Q: What do you think this query does?
- A: inserts all the actor_ids from the Film_Actor relation into the Actor relation



Deleting Tuples

- To delete tuples satisfying a condition from some relation:
- **DELETE FROM Relation WHERE <condition>;**
- Actor (actor_id, first_name, last_name, last_update)
- Q: To delete all the many “Driver”s — Adam, Brenda, Charles, and Didis from our Actor relation, what query would we use?



Deleting Tuples

- To delete tuples satisfying a condition from some relation:
- **DELETE FROM Relation WHERE <condition>;**
- Actor (actor_id, first_name, last_name, last_update)
- Q: To delete all the many “Driver”s — Adam, Brenda, Charles, and Didis from our Actor relation, what query would we use?
- **DELETE FROM Actor WHERE last_name = 'Driver';**
- To delete all tuples from a relation simply say:
- **DELETE FROM Actor;**



Updates

- To change certain attributes of certain tuples in a relation:

UPDATE Relation

SET <list of attribute assignments>

WHERE <condition>



Example: Update

- Actor (actor_id, first_name, last_name, last_update)
- Change Adam Driver's actor_id to 6000:

UPDATE Actor

SET actor_id = 6000

WHERE first_name = 'Adam' AND last_name = 'Driver'



A Tricky Issue of Semantics

- Now, we know at a high level how **INSERTS**, **DELETES**, and **UPDATES** work
- There's a special issue that impacts both **DELETES** and **UPDATES**
- Suppose we have a relation Employee (Name, Salary, Address, Title)
- We run the following update to increase salaries of all database designers who have a Salary < 100k by 20%
 - Q: How would we do this?



A Tricky Issue of Semantics

- Now, we know at a high level how **INSERTS**, **DELETES**, and **UPDATES** work
- There's a special issue that impacts both **DELETES** and **UPDATES**
- Suppose we have a relation Employee (Name, Salary, Address, Title)
- We run the following update to increase salaries of all database designers who have a Salary < 100k by 20%

UPDATE Employee SET Salary = Salary * 1.2

WHERE Title= 'Database Designer' AND Salary <100000



A Tricky Issue of Semantics

UPDATE Employee SET Salary = Salary * 1.2

WHERE Title = 'Database Designer' AND Salary < 100000

- Suppose the Employee relation has a tuple corresponding to a Database Designer “Bob” with Salary < 100000.
- When the **UPDATE** statement hits Bob’s tuple, it changes that Salary to 120% of the original one.
- If we’re not careful, we could keep increasing Bob’s Salary...
- Q: when will we stop?



A Tricky Issue of Semantics

UPDATE Employee SET Salary = Salary * 1.2

WHERE Title = 'Database Designer' AND Salary < 100000

- Suppose the Employee relation has a tuple corresponding to a Database Designer “Bob” with Salary < 100000.
- When the **UPDATE** statement hits Bob’s tuple, it changes that Salary to 120% of the original one.
- If we’re not careful, we could keep increasing Bob’s Salary...
- Q: when will we stop?
- We will stop when the Salary hits 100000 or above.



A Tricky Issue of Semantics

UPDATE Employee SET Salary = Salary * 1.2

WHERE Title = 'Database Designer' AND Salary < 100000

- How do we deal with this issue?
- We do **UPDATES** (and **DELETES**) in two phases:
 - Phase 1:
 - We look at all the tuples and evaluate their **WHERE** condition
 - If the tuple matches the **WHERE** condition, we mark it for an update or delete
 - Phase 2:
 - We actually update/delete all the marked tuples
- Homework: Convince yourselves that this always returns the same answer independent of the order of processing of tuples



Quick Demo

- Insert into the Actor relation

```
INSERT INTO Actor VALUES (5000, 'Adam', 'Driver', 'Feb 3, 2020');
```

```
INSERT INTO Actor VALUES (5001, 'Brenda', 'Driver', 'Feb 3, 2020'), (5002, 'Charles', 'Driver', 'Feb 3, 2020');
```

```
INSERT INTO Actor (actor_id, first_name, last_name, last_update) VALUES (5003, 'Didi', 'Driver', 'Feb 3, 2020');
```

- Updates to the Actor relation

```
UPDATE Actor
```

```
SET actor_id = 6000
```

```
WHERE first_name = 'Adam' AND last_name = 'Driver'
```

- Deletes from the Actor relation

```
DELETE FROM Actor WHERE last_name = 'Driver';
```



Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL Queries:
 - SFV and its semantics
 - LIKE, AS, *, %, ...
 - Null values
 - Single and multiple relations
 - Subqueries
 - Bag algebra and set operations
 - Grouping
 - Ordering
 - Case
- SQL Modifications

