

# Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL Queries
- SQL Modifications
- SQL DDL
- Next: Database Design



# So far...

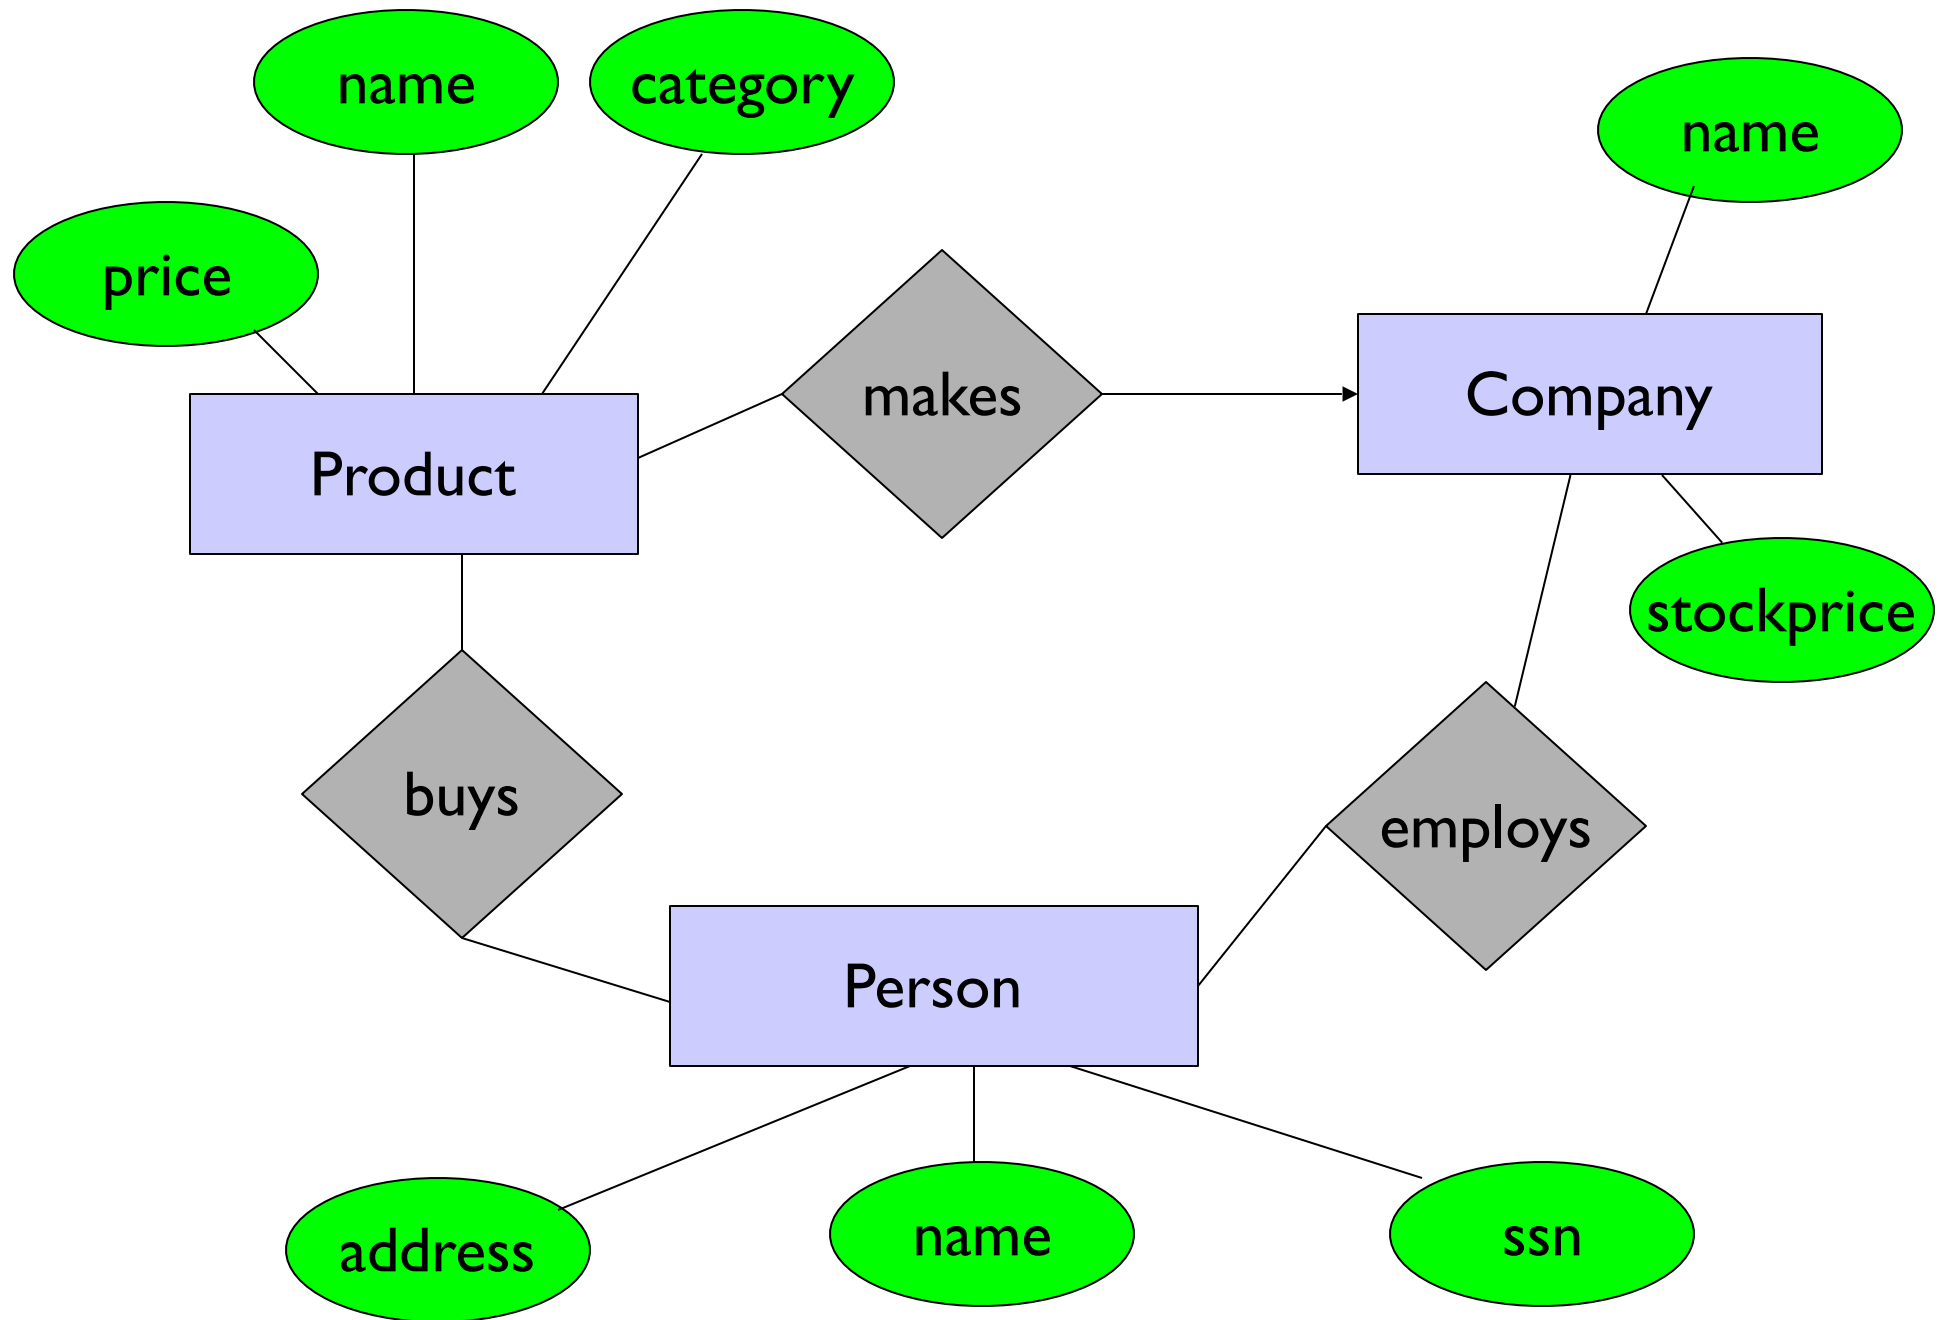
- We've seen the DDL syntax of SQL to define relations, alter them, and then drop them
- But we haven't figured out which relations to use in the first place!
- That is the focus of this set of slides



# To Decide on the Relational Schema...

- You need to faithfully capture the requirements of the application
  - What information needs to be represented and how they relate with each other
- A step towards that is a conceptual model, or a diagram of the entities and relationships
  - This is the so-called ER diagram or Entity-Relationship diagram
- Essentially “doodles about data”




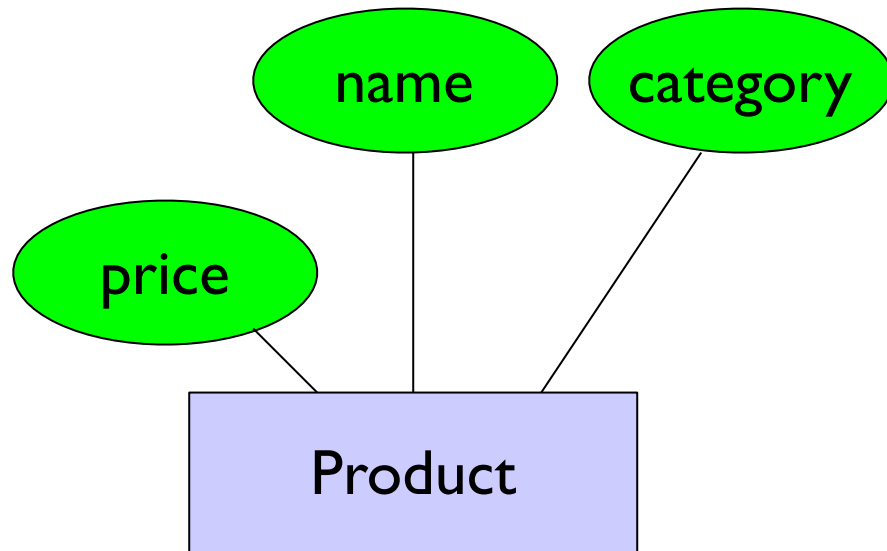


# The ER diagram

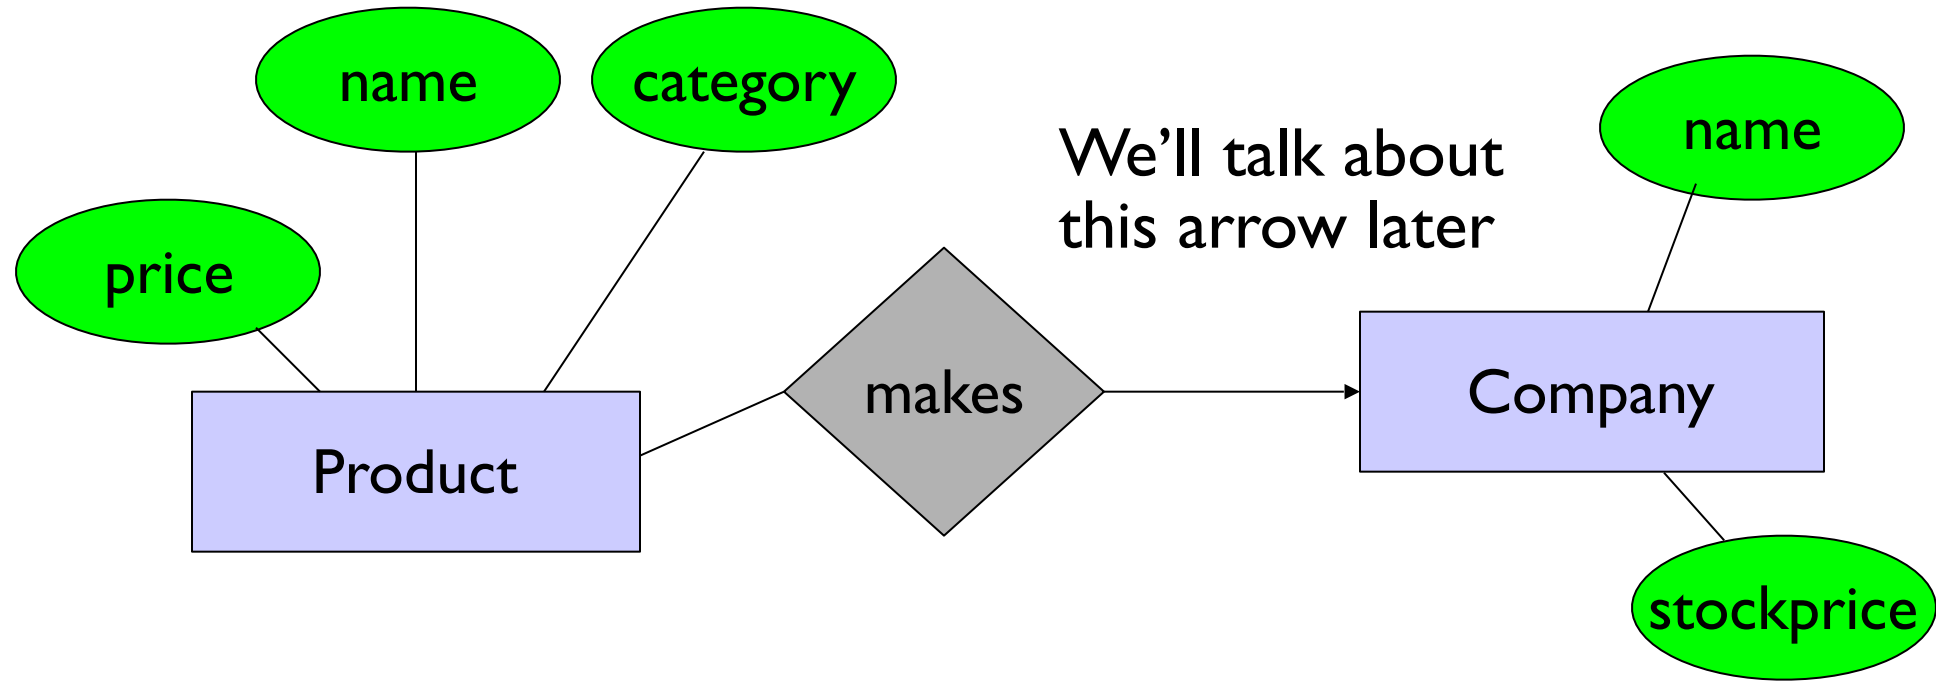
- Proposed by Peter Chen in 1976
- A very rich modeling language: we're going to study a tiny subset
- Two basic primitives in an ER diagram:
  - **Entities**: things, objects, ...
    - Described using a set of attributes (all atomic)
    - A set of entities is called an **Entity set**
      - Represented via a rectangle
  - **Relationships**: connections between entities
    - Represented via diamonds



# Entity Sets



# Relationships

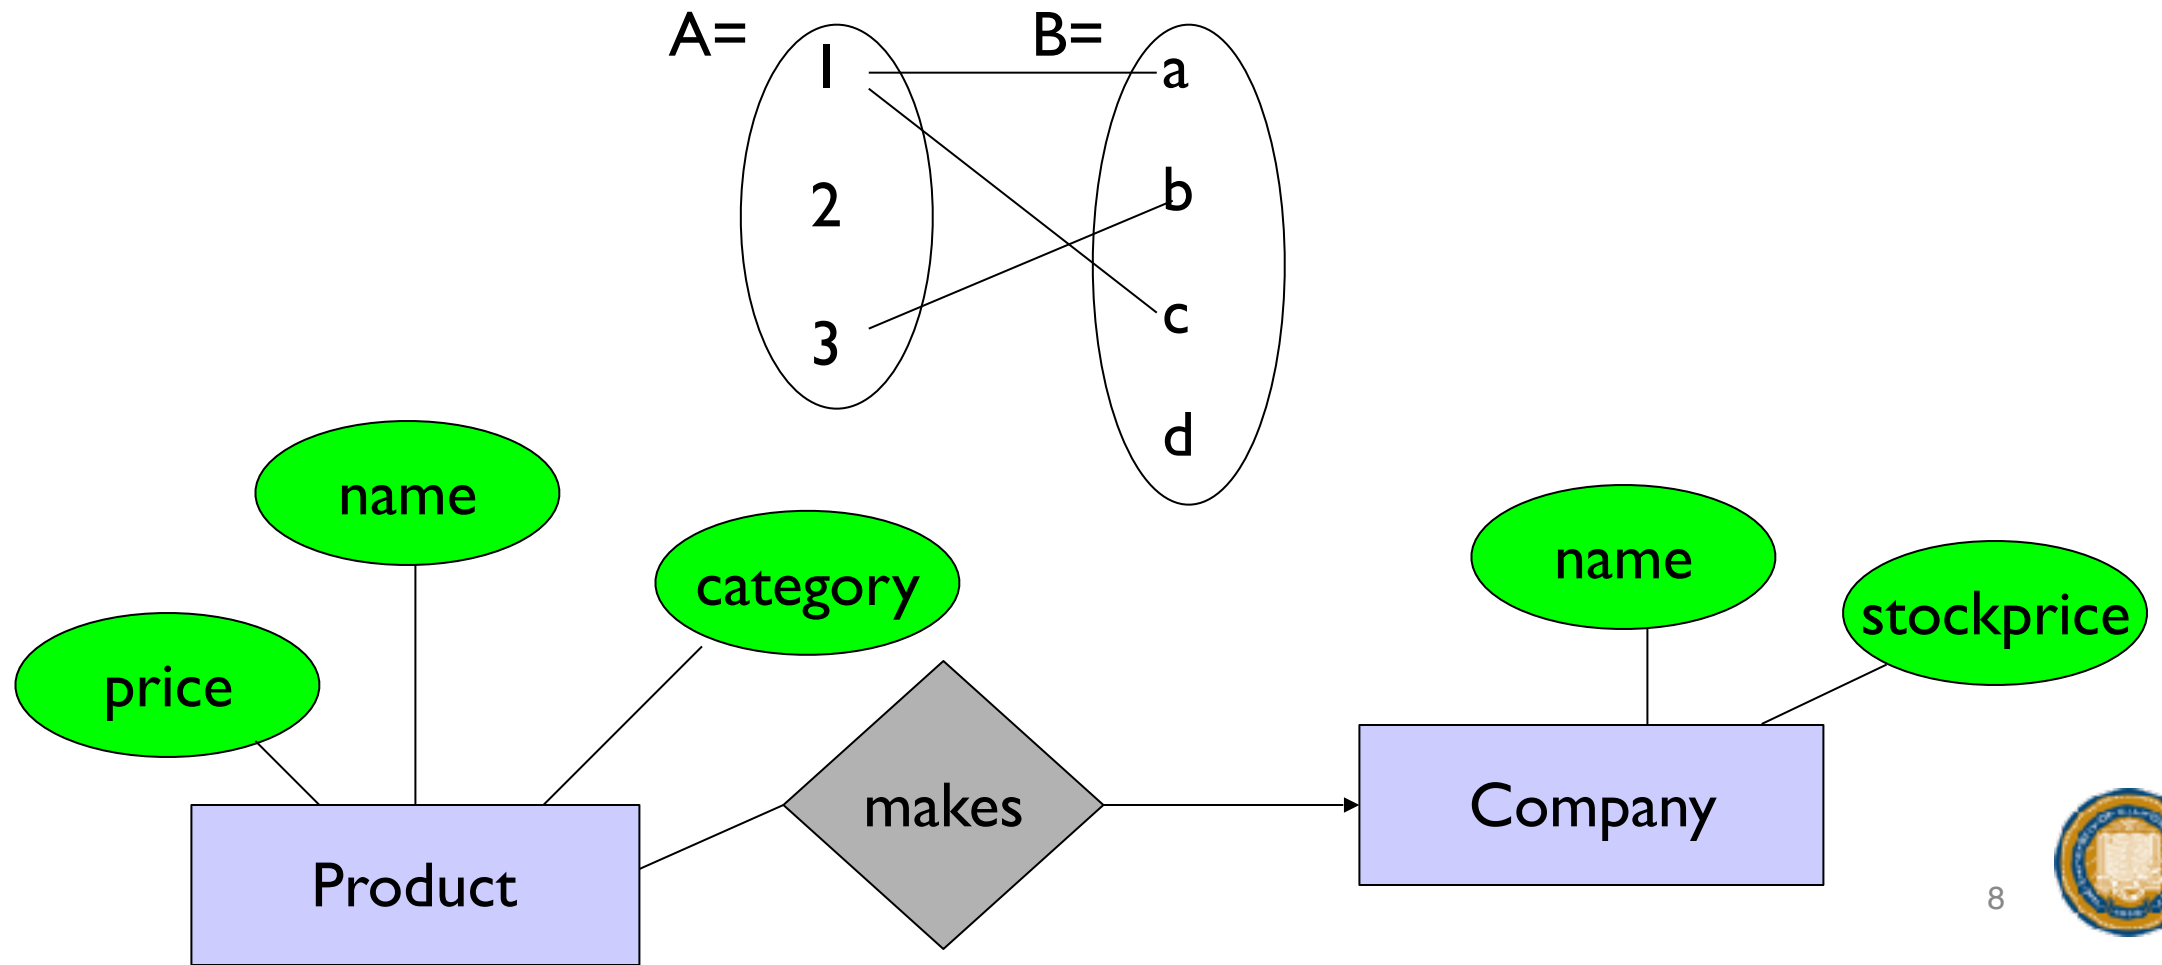


Goal: Connect Entity Sets



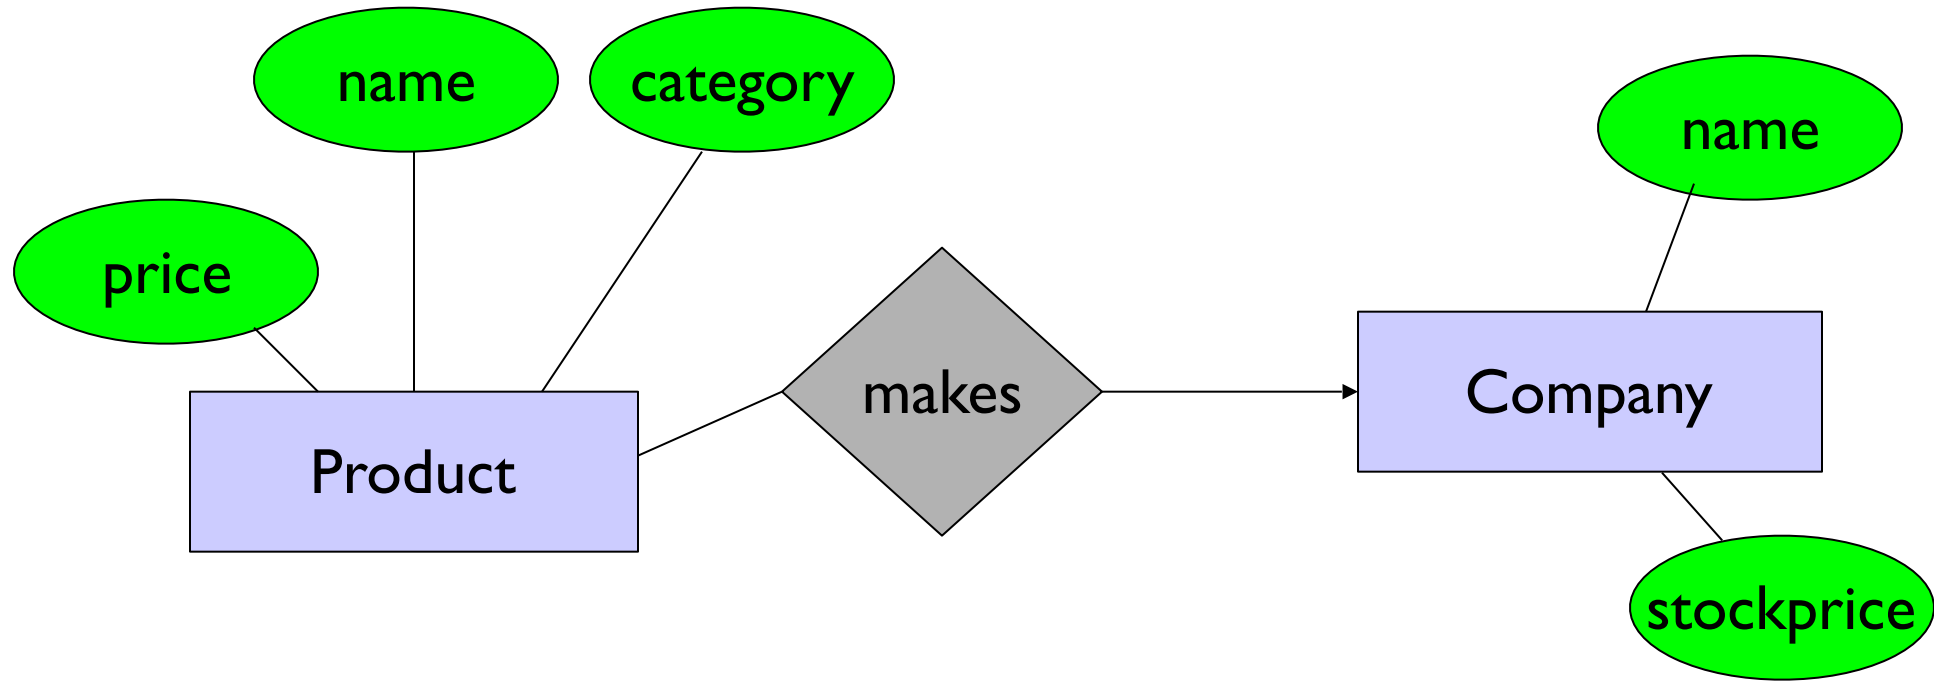
# Relationships

- A and B are Entity Sets, a Relationship between A and B is a subset of  $A \times B$
- Makes is a subset of Product  $\times$  Company





# Back to the Diagram






- Entity Sets: Rectangles
- Attributes: Ovals attached to rectangles or diamonds
- Relationships: Diamonds connecting rectangles






# The Arrows in Relationships

- Arrow = **at most one**
- Another interpretation: *determines*
- One-one: one on LHS/RHS connected to at most one of RHS/LHS
- Many-one: one on LHS connected to at most one on RHS
- Many-many: no constraints



# Example Scenarios for Each Case

- one-one:
- many-one
- many-many



Department – head ??




Actor – play ??

Employee – company ??



# Example Scenarios for Each Case

- one-one:
- many-one
- many-many



Department – head:

one – one. (each dep has 1 head, each head is head of 1 dep)


Actor – play: many – many. (each play has many actors, and vice versa)

Employee – company:

many – one. (each company has many employees, each employee works for one company)



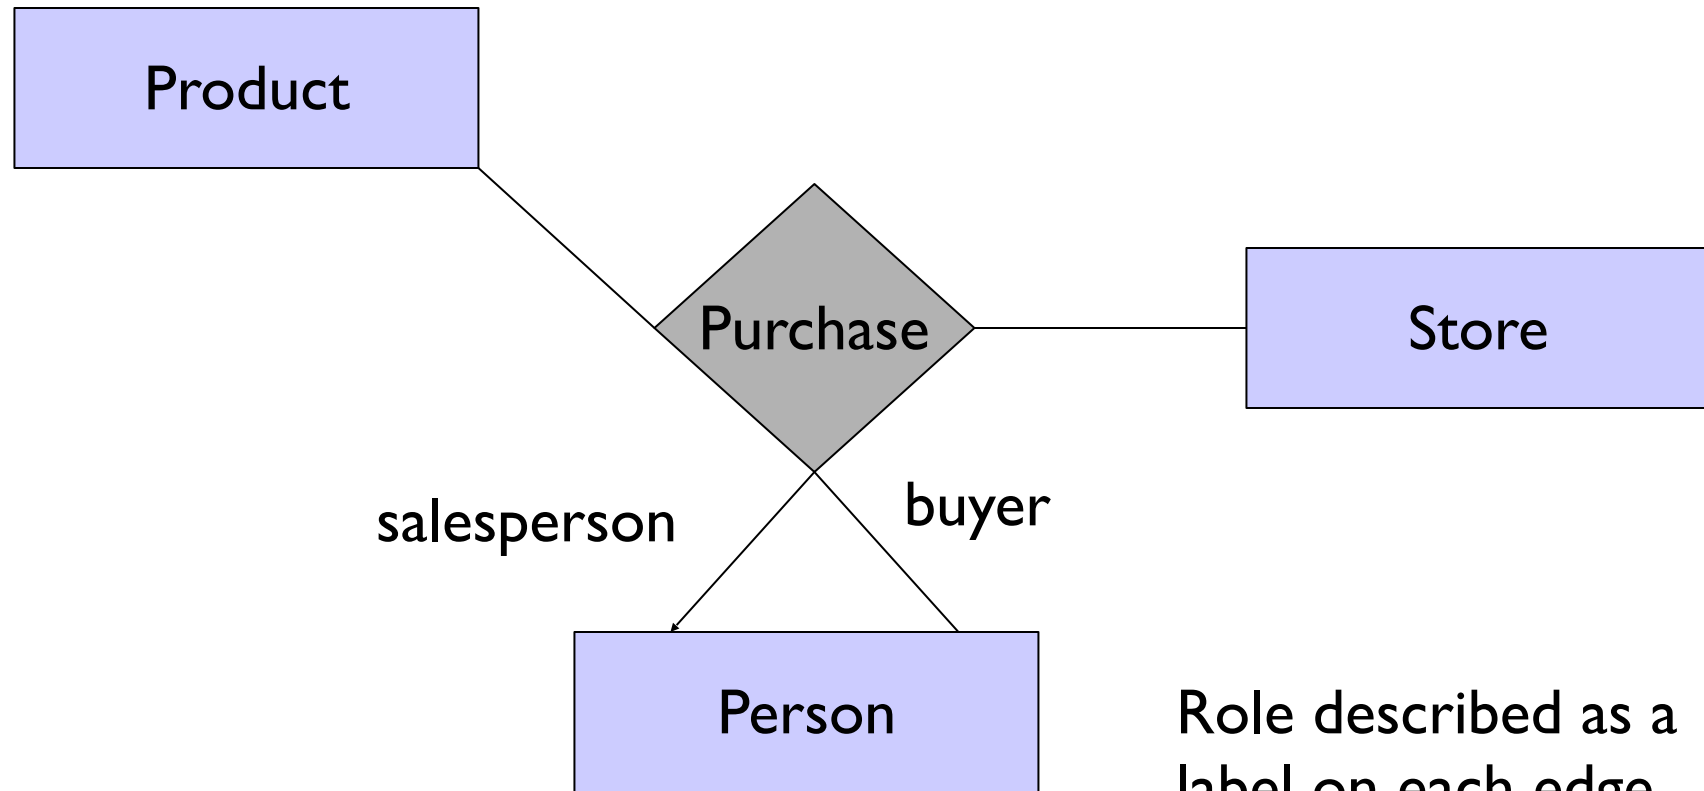
# Multi-way Relationships



- Relationship between person, store, purchase
  - Purchase is a subset of  $\text{Product} \times \text{Person} \times \text{Store}$
- Arrow still means *at most one*:
  - **combination of the other entities determines or is connected to at most one entity of inbound entity set**



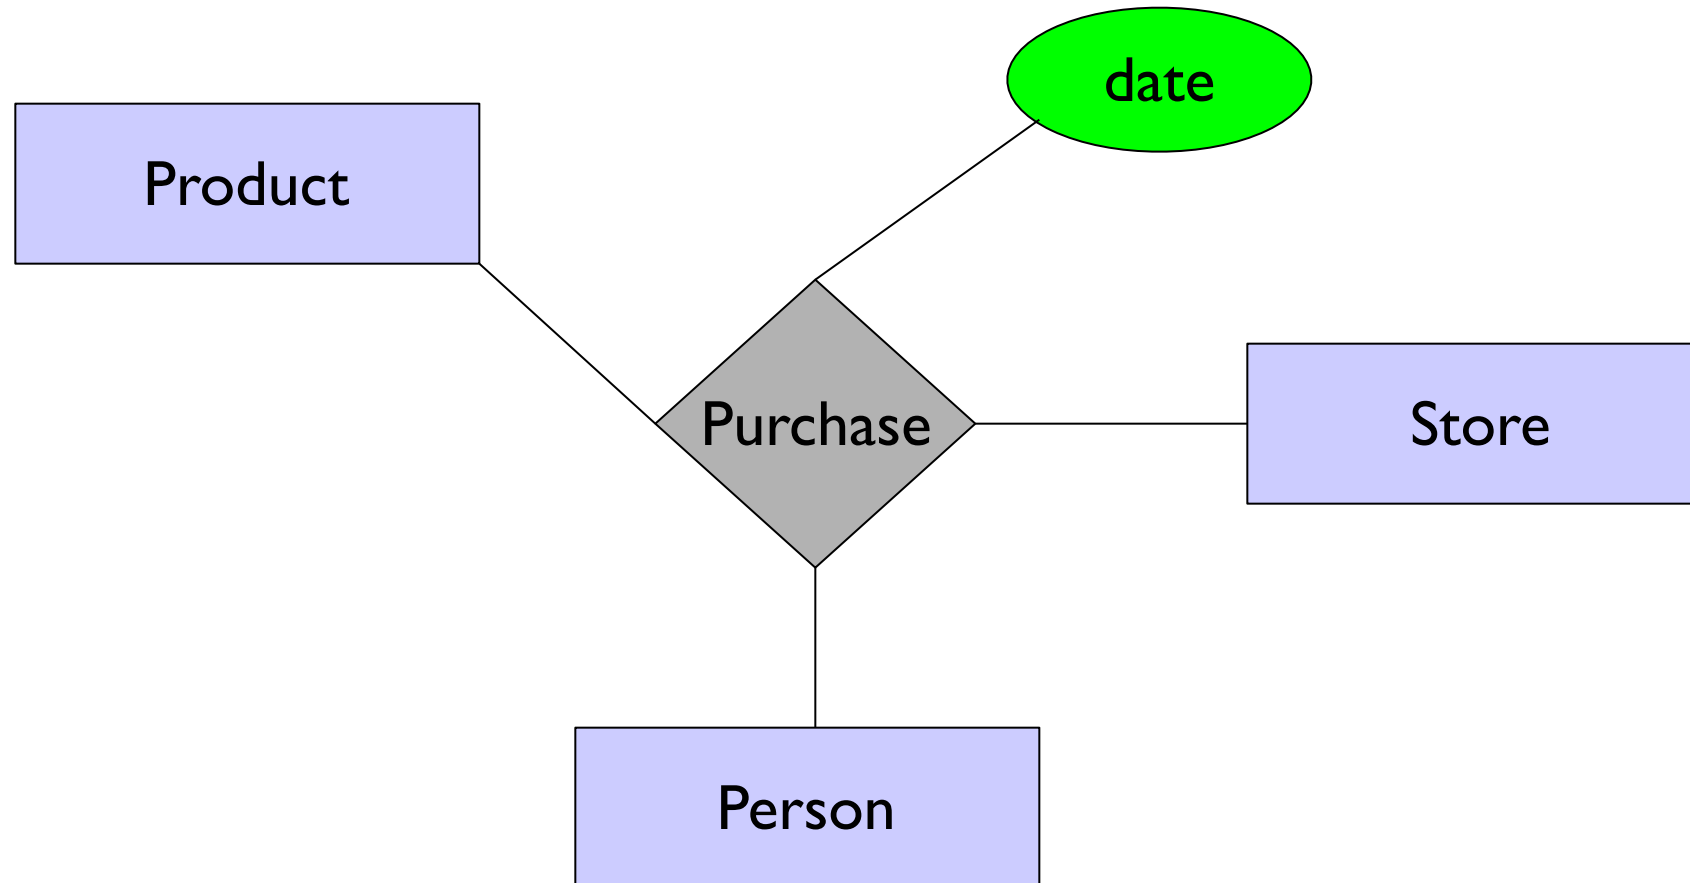
# Roles in Relationships



Role described as a label on each edge.

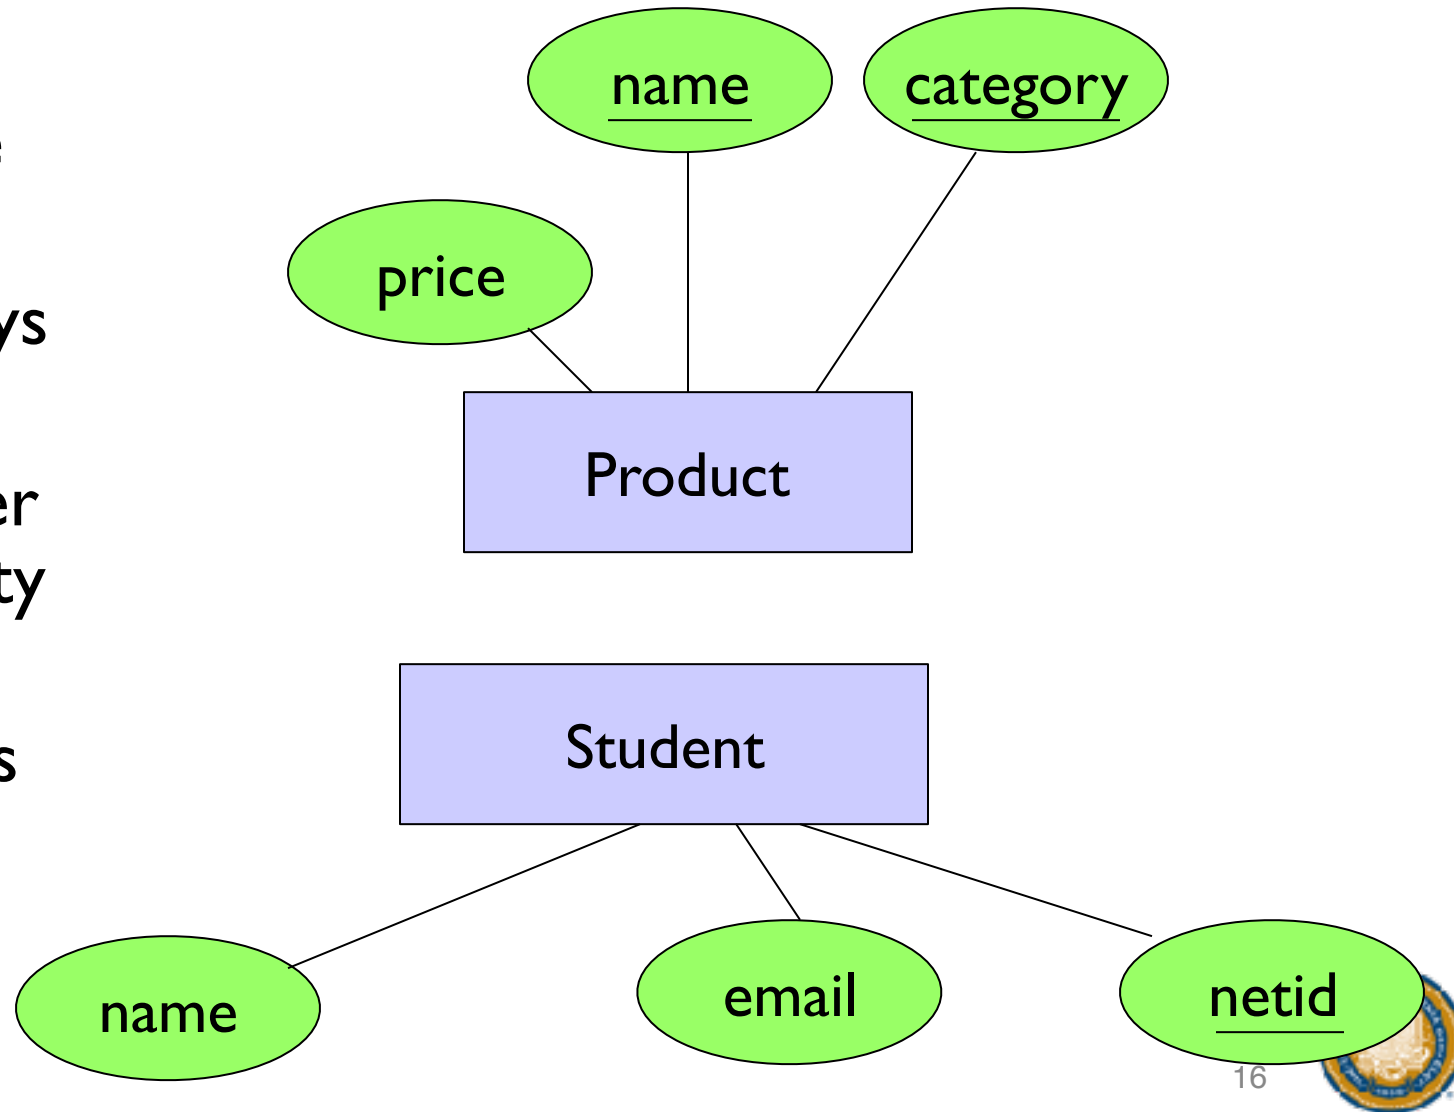


# Attributes of Relationships



# Remember Primary Keys?

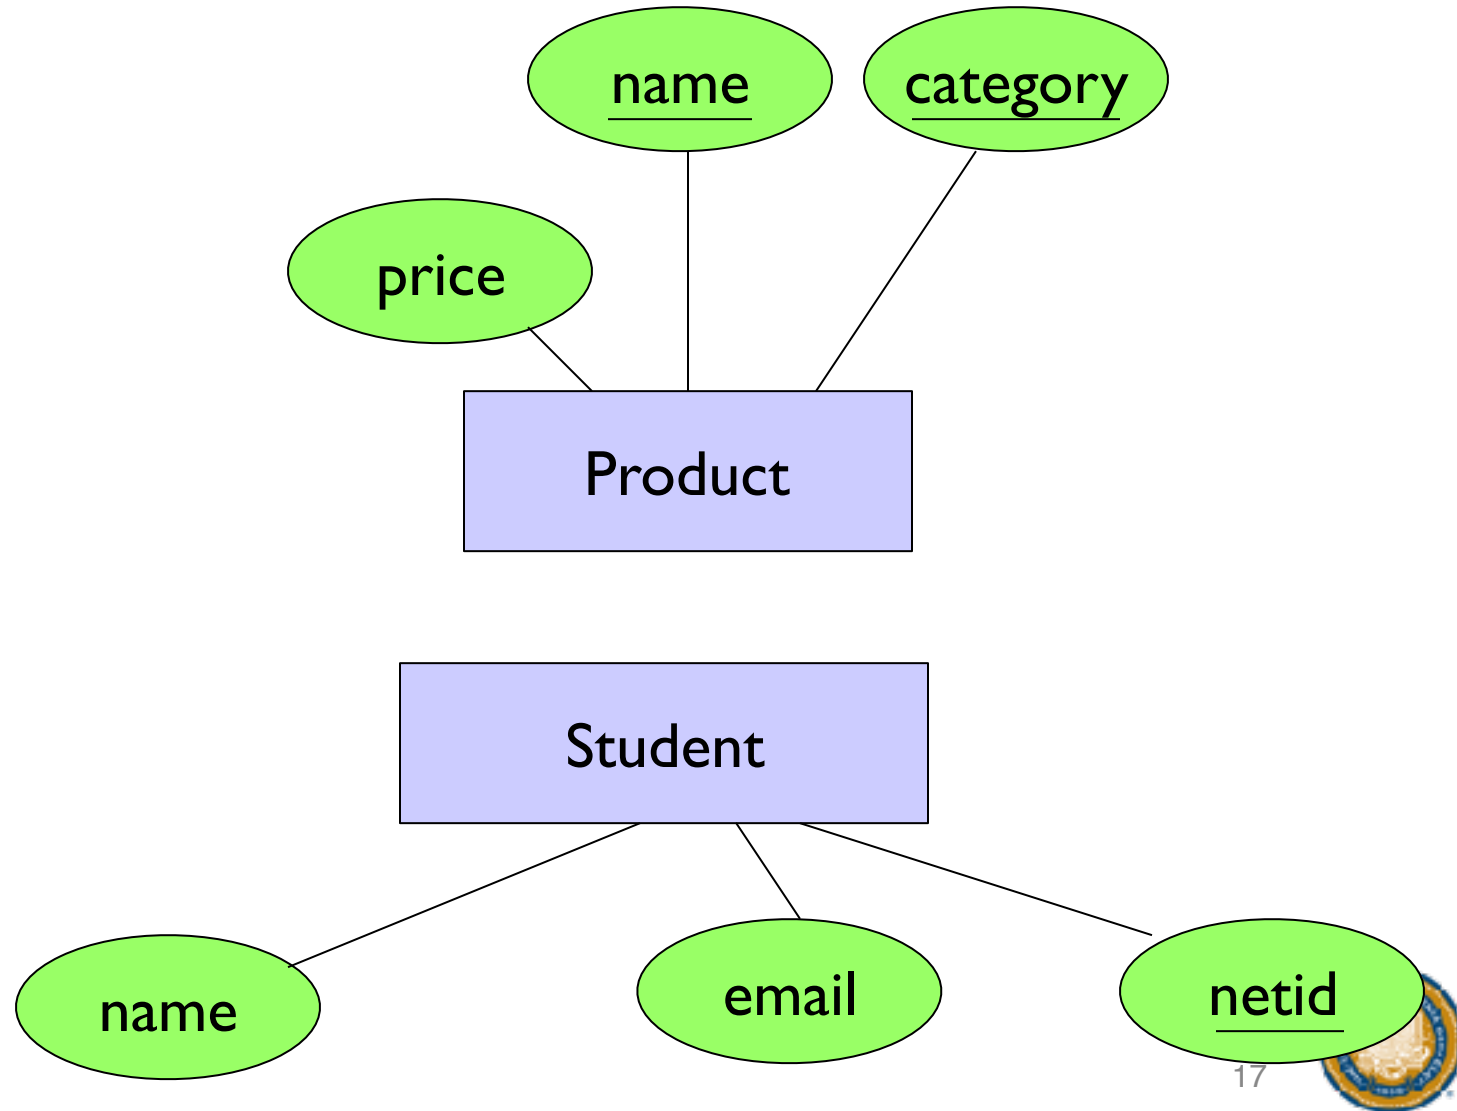
- A form of constraint on the data
- We can capture primary keys by underlining them
- Name and category together determine the Product entity
- No formal way to capture multiple keys in ER diagrams






# Remember Primary Keys?

- In ER diagrams, we require every entity set to have a primary key
- In the worst case it is the entire set of attributes!




# Another Constraint: Referential Integrity

- Recall: arrow meant “at most one”
- Each product made by (related to) at most one company




- But wouldn't it be weird if a product wasn't made by any company?



- Encode that with a semi-circle: **exactly one**, as opposed to arrow: **at most one**
- **“No dangling pointers”**



# Another Example



What do these two semi-circles mean?

Each movie is owned by precisely one studio, and  
Each president runs exactly one studio

Each studio has up to one president



# Lots More!

There's a lot more that ER diagrams can capture...

- Subclasses
- “Weak” entity sets
- Other constraints
- ...




# Design Principles


- Be faithful to reality
- Avoid redundancy
- Pick the right kind of element




# Being Faithful to Reality



No: A Person may purchase multiple Products



No: A Country can have at most one President

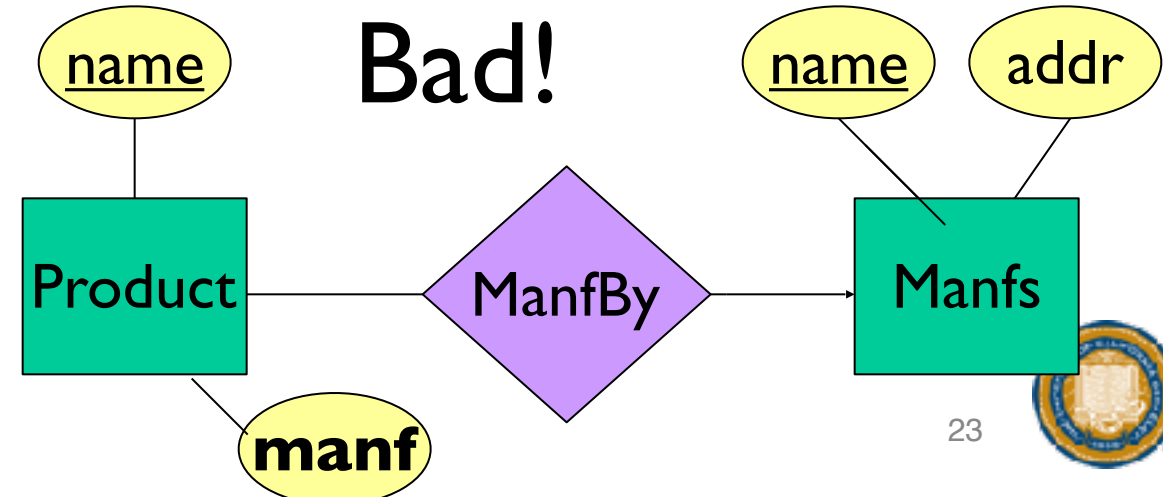
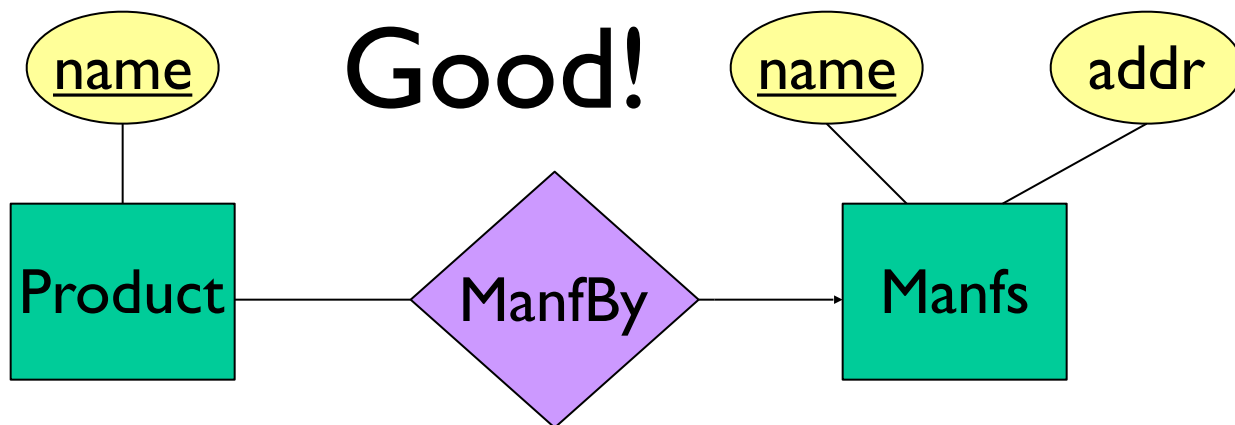


Yes if multiple instructors, No if not.



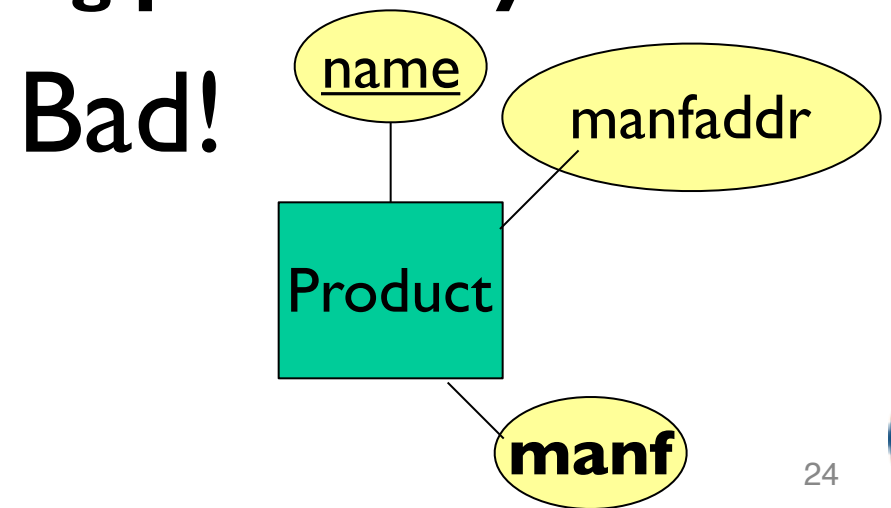
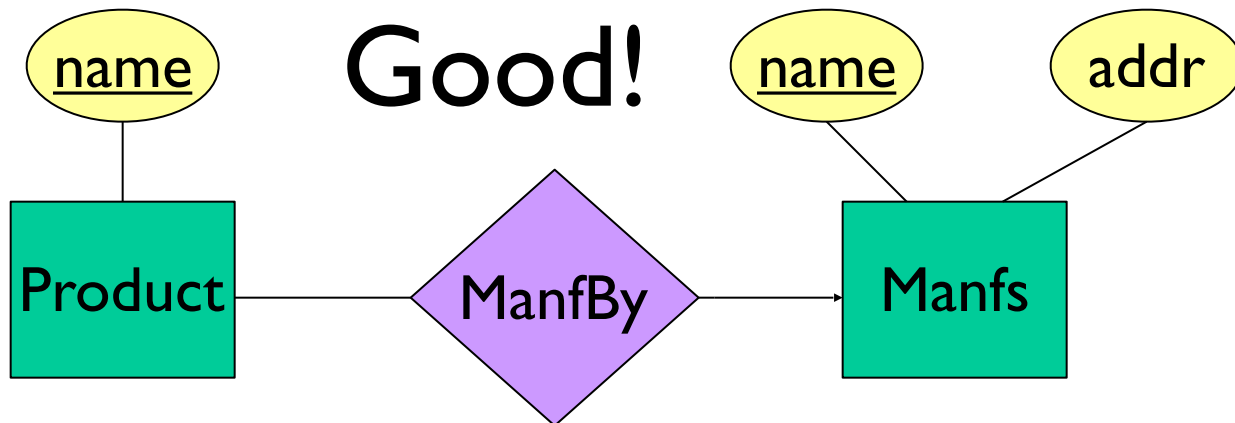
# Avoiding Redundancy

- Redundancy: saying the same thing in multiple ways or places
- Redundancy
  - Wastes space
  - Allows inconsistencies to pop up (**update anomalies**)
    - If you edit one copy but not others



# Avoiding Redundancy


- Redundancy: saying the same thing in multiple ways or places
- Redundancy
  - Wastes space
  - Allows inconsistencies to pop up
  - **Another issue with this one: you lose addresses for manf who haven't started making products yet**





# Picking the Right Kind of Element

- An Entity Set should satisfy one of the following conditions:
  - It is more than just the name, i.e., it has at least one non-key attrib
  - OR, it is the “many” in a many-one or many-many relationship




- Product: “many” of a many-one relationship
- Manf: has a non-key attrib: addr



# Picking the Right Kind of Element

- An Entity Set should satisfy one of the following conditions:
  - It is more than just the name, i.e., it has at least one non-key attrib
  - OR, it is the “many” in a many-one or many-many relationship



- If we had no manf address info, might as well merge into product



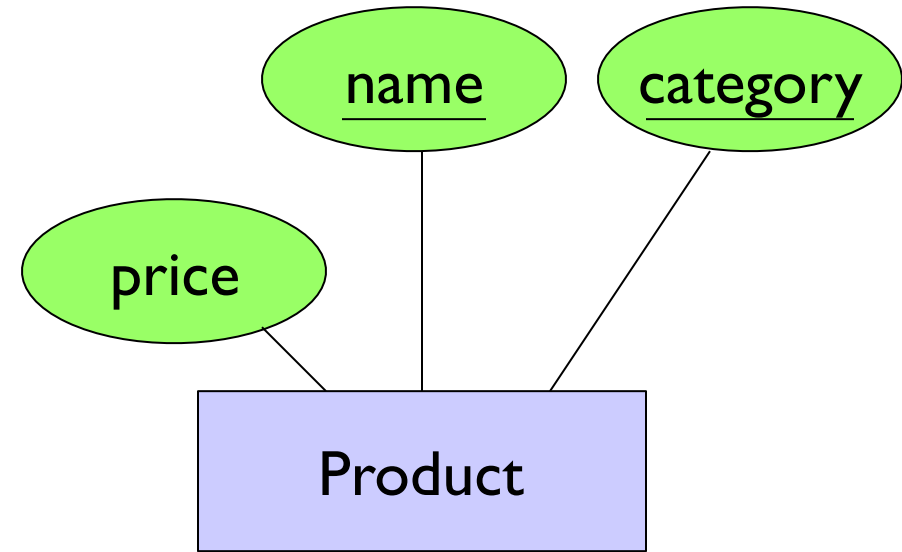
# Next: Translation into a Relational Schema

- Now, let's say we have an ER diagram
- How do we convert it into a relational schema?
- Turns out many database (ER) design software will do this automatically for you!



# Entity Set to Relation

- Relation:  
Product (name, category, price)

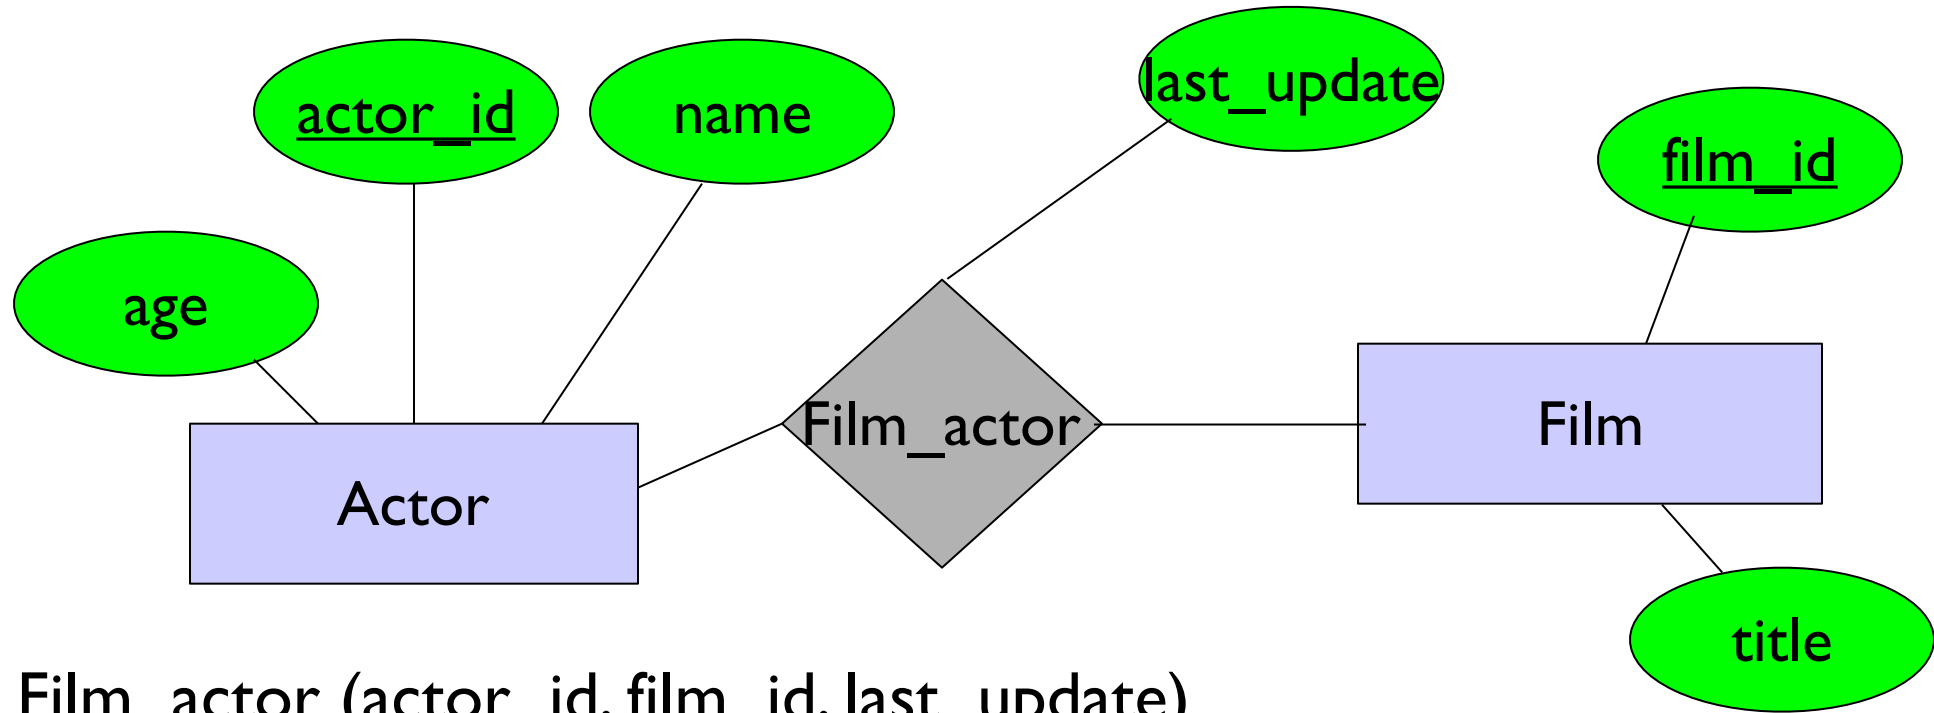


CREATE TABLE Product

(name VARCHAR (25), price REAL, category CHAR (4),  
PRIMARY KEY (name, category));



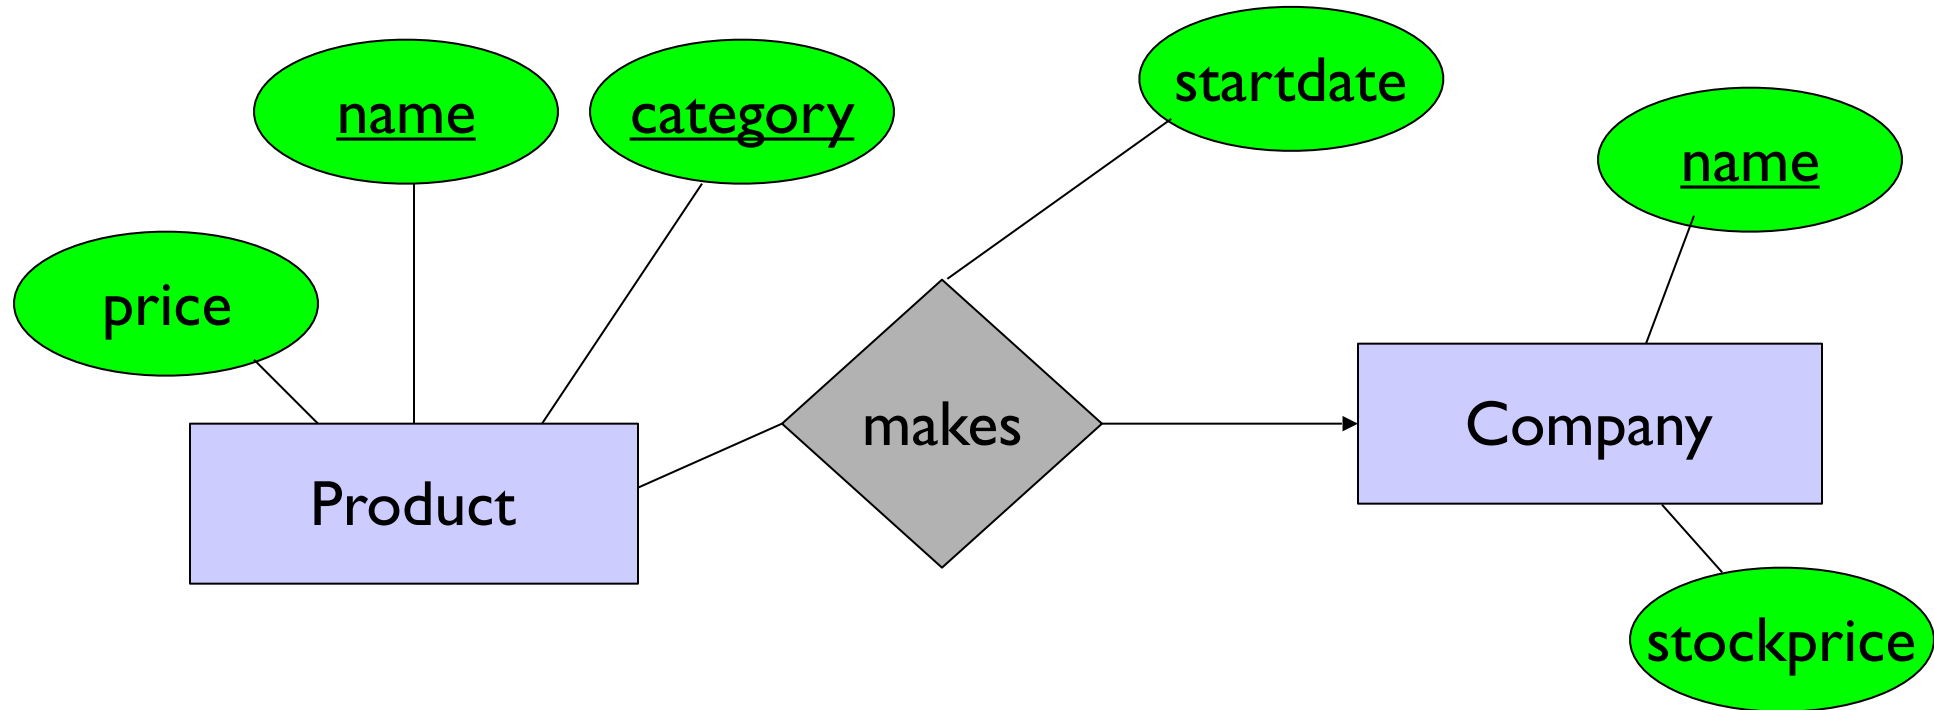
# Relationship to Relation: Our Familiar Movie Database



- Film\_actor (actor\_id, film\_id, last\_update)
- Include:
  - Key attributes of each Entity Set
  - Attributes of the relationship
- Three tables: Actor, Film\_actor, Film



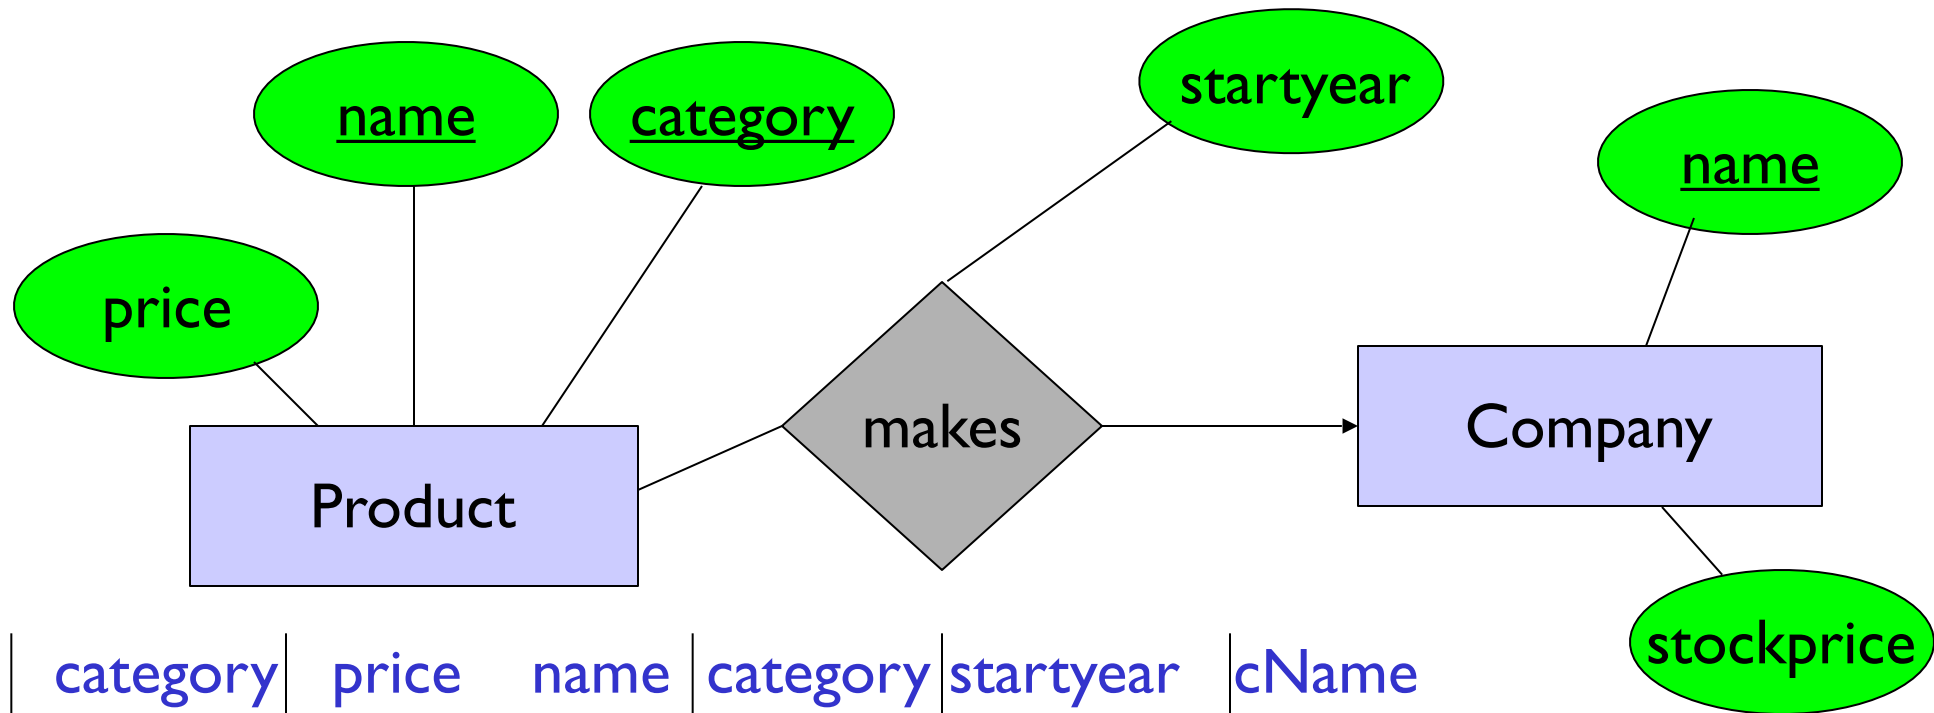
# Relationship to Relation: Many to one



- Can improve the representation if it is many-to-one (or one-to-one)
  - Product (price, name, category, startdate, companyname)
  - Instead of:
    - Product (price, name, category),
    - Makes (name, category, startdate, companyname)



# Relationship to Relation: Many to one



<u>name</u>	<u>category</u>	<u>price</u>	<u>name</u>	<u>category</u>	<u>startyear</u>	<u>cName</u>
gizmo	gadgets	19.99	gizmo	gadgets	1963	gizmoWorks
kphone	phone	200.00				

<u>name</u>	<u>category</u>	<u>price</u>	<u>startyear</u>	<u>companyName</u>
gizmo	gadgets	19.99	1963	gizmoWorks
kphone	gadgets	200.0	NULL	NULL



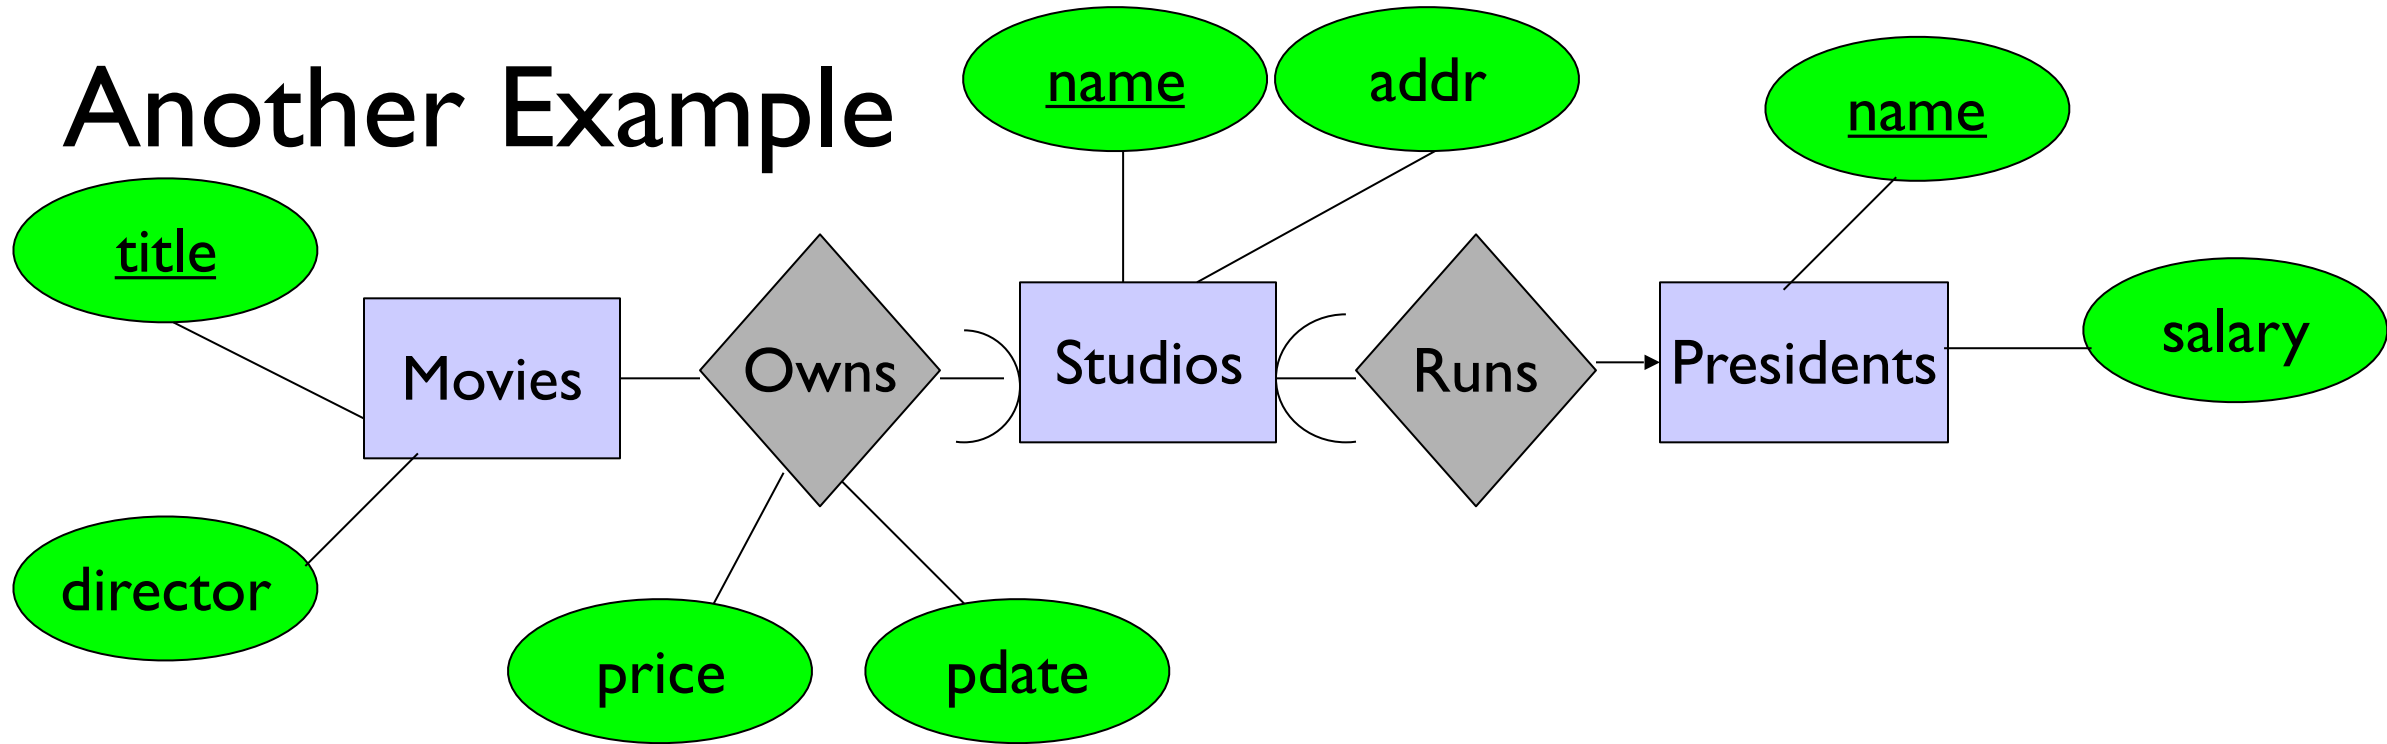
# Combining Relations

- It is OK to combine the relation for an entity-set  $E$  with the relation for  $R$  if  $R$  is a many-one relationship from  $E$  to another entity set.
- Typically, when combining two tables into one, there is always a danger of redundant information: the same information represented multiple times
  - Why does the redundancy argument not apply here?
    - One single tuple “smushed” together





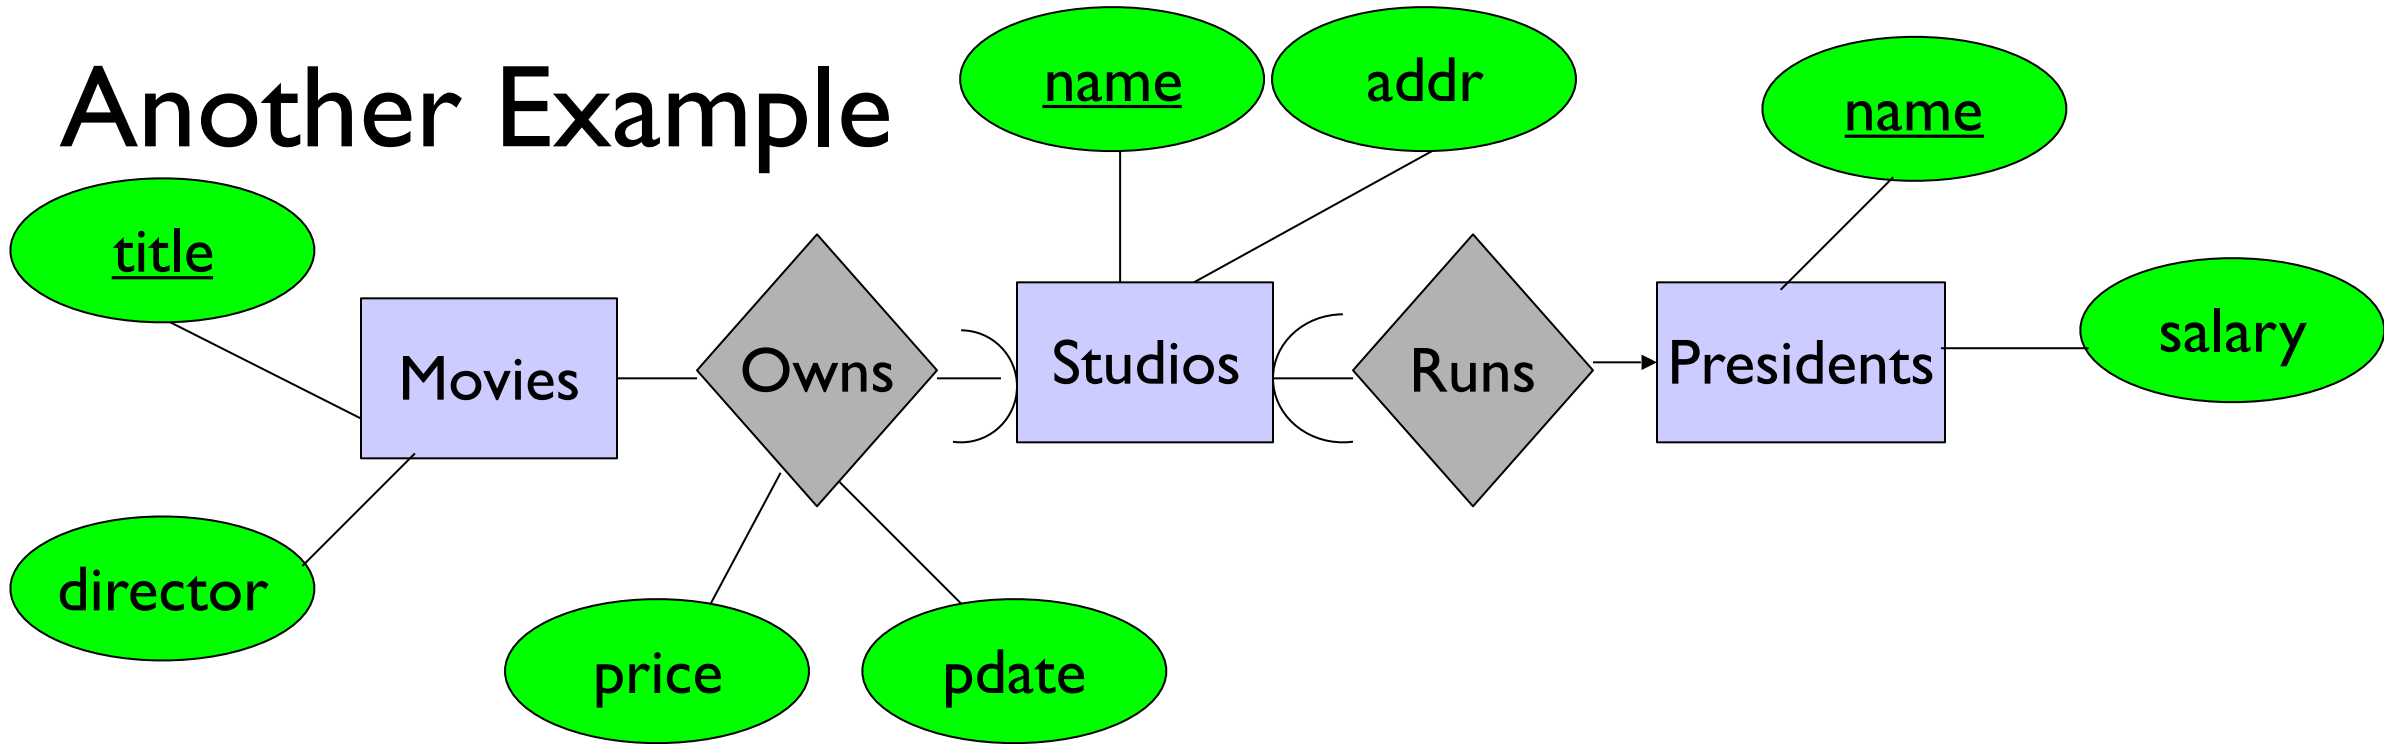
# Another Example



- What relations would we have if we had one relation for each Entity Set and Relationship?



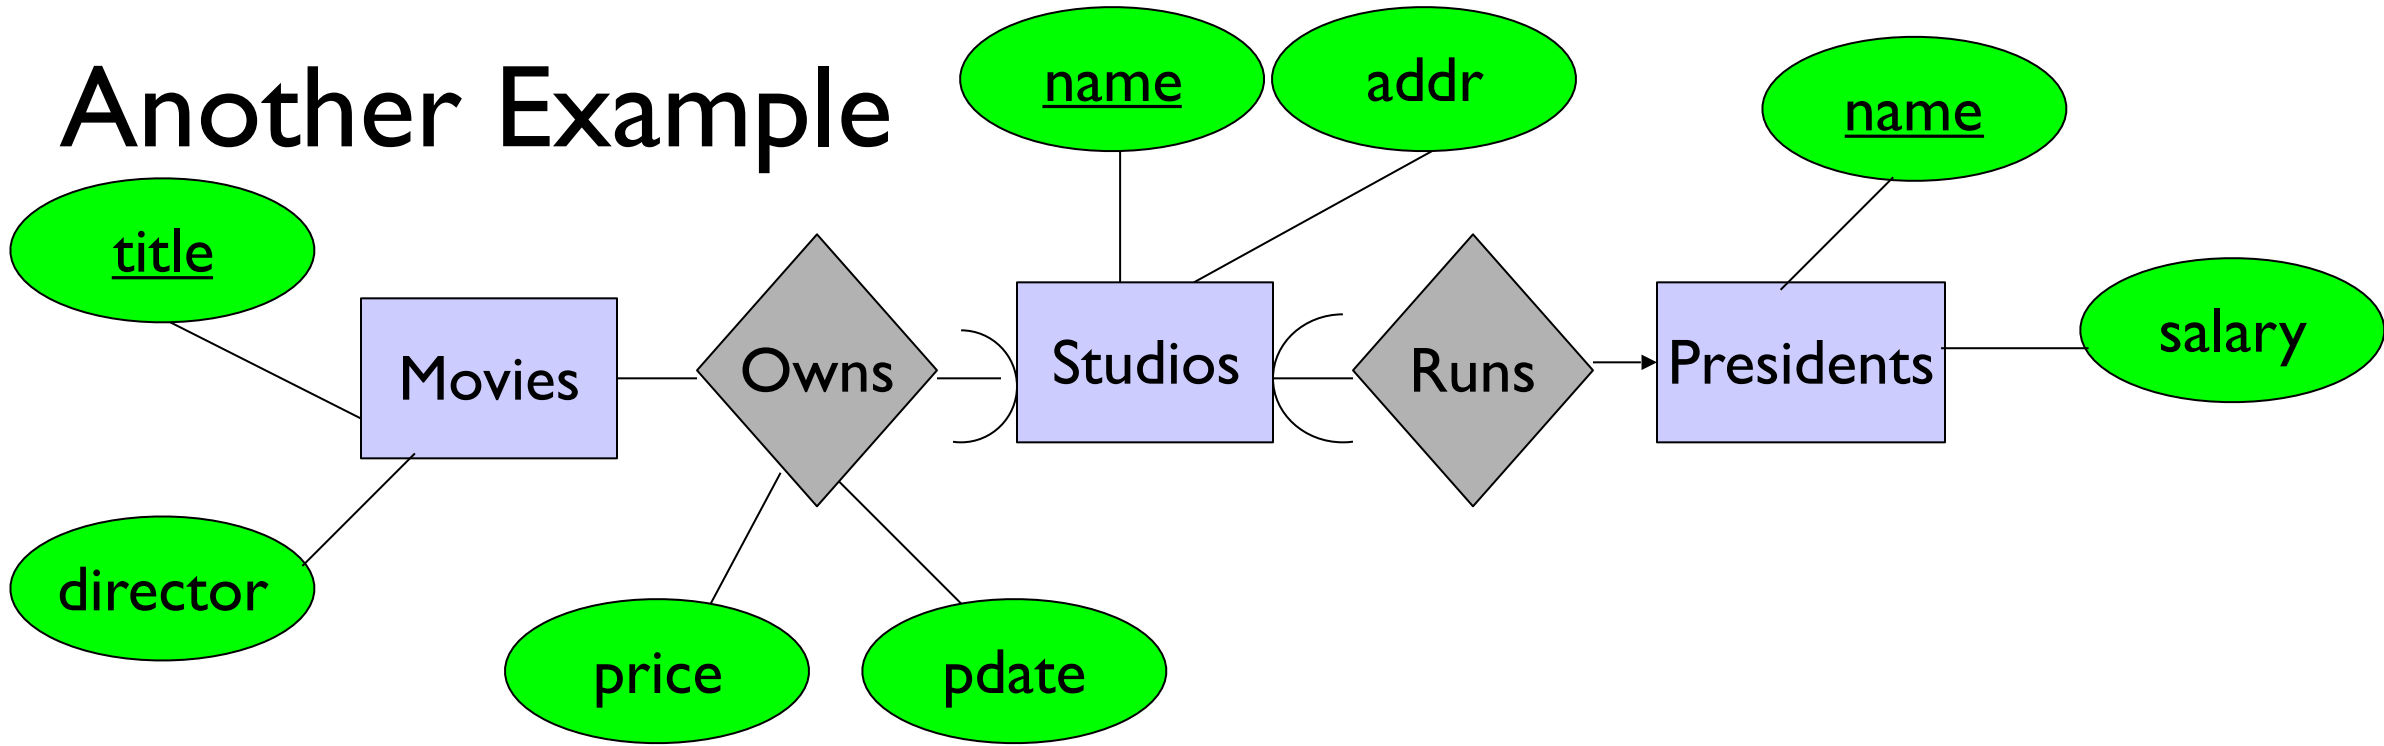
# Another Example



- Movies (title, director)
- Owns (movie\_title, price, pdate, studio\_name)
- Studios (name, addr)
- Runs (studio\_name, president\_name)
- Presidents (name, salary)




# Another Example



- Movies (title, director, **price**, **pdate**, **studio\_name**)
- ~~Owns (movie\_title, price, pdate, studio\_name)~~
- Studios (name, addr)
- Runs (studio\_name, president\_name)
- Presidents (name, salary)




# Another Example



- Movies (title, director, **price**, **pdate**, **studio\_name**)
- ~~Owns (movie\_title, price, pdate, studio\_name)~~
- Studios (name, addr, **president\_name**)
- ~~Runs (studio\_name, president\_name)~~
- Presidents (name, salary)




# Another Example



- Movies (title, director, **price, pdate, studio\_name**)
- ~~Owns (movie\_title, price, pdate, studio\_name)~~
- Studios (name, addr, **president\_name, salary**)
- ~~Runs (studio\_name, president\_name)~~
- ~~Presidents (name, salary)~~



# Another Example



- Movies (title, director, **price**, **pdate**, **studio\_name**)
- Studios (name, addr, **president\_name**, **salary**)
- Why not go all the way and have all the information in one table?
- Massive redundancy!



# One Final Piece: Functional Dependencies and Normalization

- Start with requirements, then abstract them out into an ER diagram, convert them into relations
  - Are we done?
  - **Not quite! There can still be issues with the relations thus designed, causing us problems**
  - Or, we may start with a set of relations (not so carefully designed without ER diagrams) and have to “fix it”
- The concepts of “functional dependencies” and “normalization” help us find our way to good designs in such cases
  - We’ll do a very quick intuitive overview... but will skip detailed discussions
- Let’s consider an example



# Is this a good design?

- Individuals with several phones

Address	SSN	Phone Number
10 Green	123-456-789	(201) 233-1456
10 Green	123-456-789	(201) 123-3439
431 Purple	987-654-321	(145) 241-2131
431 Purple	987-654-321	(312) 123-1287





# Is this a good design?

- Individuals with several phones
- Redundancy:
  - Address repeated multiple times
- Update anomalies:
  - If we update address of person with phone number 201-233-1456, there will be two addresses of that person (wrong!)
- Deletion anomalies:
  - If we delete the phone number for a person, we need to check if there is another phone number for said person, else may lose address info.

Address	SSN	Phone Number
10 Green	123-456-789	(201) 233-1456
10 Green	123-456-789	(201) 123-3439
431 Purple	987-654-321	(145) 241-2131
431 Purple	987-654-321	(312) 123-1287



# Better Designs Exist!

Address	SSN
10 Green	123-456-789
431 Purple	987-654-321

SSN	Phone Number
123-456-789	(201) 233-1456
123-456-789	(201) 123-3439
987-654-321	(145) 241-2131
987-654-321	(312) 123-1287

- Each bit of info only exists “once”
- How would you recover the original relation via these two relations?
- A: a natural join
- Unfortunately, will not detect this even with principled ER design and translation



# Better Designs Exist!

Address	SSN
10 Green	123-456-789
431 Purple	987-654-321

SSN	Phone Number
123-456-789	(201) 233-1456
123-456-789	(201) 123-3439
987-654-321	(145) 241-2131
987-654-321	(312) 123-1287

- Each bit of info only exists “once”
- Splitting relations into multiple relations that minimizes redundancy is called **normalization**
  - We use functional dependencies to guide us
- There can be other objectives (in addition to minimizing redundancy)



# Better Designs Exist!

Address	SSN
10 Green	123-456-789
431 Purple	987-654-321

SSN	Phone Number
123-456-789	(201) 233-1456
123-456-789	(201) 123-3439
987-654-321	(145) 241-2131
987-654-321	(312) 123-1287

- Each bit of info only exists “once”
- Benefits of normalization:
  - Removing redundancy, minimizes update, and delete anomalies
- Downsides of normalization:
  - Joins are costly sometimes; want to avoid them
  - This is the main reason why people avoid normalization



# Functional Dependencies

- How do perform normalization? use **functional dependencies!**
- A form of constraint on your data

$$A_1, A_2, \dots, A_n \rightarrow B_1, B_2, \dots, B_m$$

- If two tuples agree on values of the LHS, they must agree on the values of the RHS: **holds for all instances!**
- Q: Have we discussed any examples of FDs so far?
- Special case: LHS is the primary key (or any key)
  - Movies (title, director, price, pdate, studio\_name)
    - title  $\rightarrow$  director, price, pdate, studio\_name



# Functional Dependencies improve Designs

- Functional dependencies explicitly capture almost (see MVDs!) all interesting redundancy-oriented constraints in your data
- Provided by the user as “semantics”
- Returning to our example
  - $SSN \rightarrow Address$  [not just here]
  - $SSN \nrightarrow Phone\ Number$
- This causes us to “factor out” SSN and Address into a separate relation
  - “decompose” or “normalize”

Address	SSN	Phone Number
10 Green	123-456-789	(201) 233-1456
10 Green	123-456-789	(201) 123-3439
431 Purple	987-654-321	(145) 241-2131
431 Purple	987-654-321	(312) 123-1287



# Lots more to say about FDs and Normalization!

- Various types of “normal forms”: BCNF, 1NF, 2NF, 3NF & trade offs!
- Calculus of FDs: closure, splitting/combining, basis
- But we won't spend more time on this
- Rule: “decompose/normalize” your relations to minimize redundancy, while ensuring that you can correctly reconstruct the original relation.
  - But beware that normalization can lead to performance overheads due to joins!

