

# Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL Queries
- SQL Modifications
- SQL DDL
- Database Design
- Views
- Constraints and triggers
- Next: indexes



## Index

[illegible]

# Indexes

- Indexes help you look for data matching certain characteristics
  - Much like an index in a book! Or a library!
- An index on an attribute  $A$  allows you to find tuples with a specific value for  $A$ 
  - Think of indexes as data structures or dictionaries storing “index key-value” pairs
    - Index key = values of  $A$
    - Value = locations of tuples
    - Q: what is the correspondence for an index in a book?
  - This may be useful, for example, to find tuples with
    - $A = a$
    - $A \leq a$
    - $a \leq A \leq b$



# Why Indexes?

- Often your relations are pretty large
  - millions, billions, trillions of tuples
- Scanning through all the tuples to find those that match a certain condition can be time-consuming
  - `SELECT * FROM Film WHERE rental_duration = 5 AND rental_rate = 2.99;`
    - Naive approach: Go through all 10000 tuples in Film, verify if rental\_duration and rental\_rate are 5 and 2.99
    - Or: with an index on rental\_duration: go through 100 tuples with rental\_duration 5
    - Or: with a multi-dimensional index on rental\_duration and rental\_rate, go through 10 tuples with values 5 and 2.99 respectively



# Also Useful for Joins!

```
SELECT * FROM Film, FilmActor  
WHERE Film.film_id = FilmActor.film_id;
```

- Say Film has 10,000 tuples and FilmActor has 50,000 tuples
  - Roughly 5 actors per film (say)
- Q: If you have no indexes, need to consider how many pairs of tuples?
  - $10,000 * 50,000 = 5 * 10^8$
- Q: How would you use indexes here?
  - For each film in Film, look up tuples in FilmActor, or vice versa
  - Goes from  $5 * 10^8$  to  $5 * 10^4$  tuples considered



# Also Useful for Joins!

```
SELECT * FROM Film, FilmActor  
WHERE Film.film_id = FilmActor.film_id  
AND rental_duration = 5 AND rental_rate = 2.99
```

- Say we had indexes for FilmActor.film\_id, Film.film\_id, and Film.rental\_duration. Q: How would you use indexes here?
  - Lookup films with rental\_duration 5 = 100 tuples
  - For each one, look up FilmActor = 100 \* 5 tuples matched
  - Verify if rental\_rate = 2.99



# Declaring Indexes

`CREATE INDEX rentIndex ON Film (rental_duration)`

- Allow for lookups based on rental\_duration

`CREATE INDEX rentReleaseIndex ON Film (rental_duration, release_year)`

- Imagine this to be a “index key-value” pair dictionary based on concatenating attribute values
- Allow for lookups based on pairs of rental\_duration and release\_year
- Or just for rental\_duration
- Q: Why not release\_year only?

`DROP INDEX rentReleaseIndex`



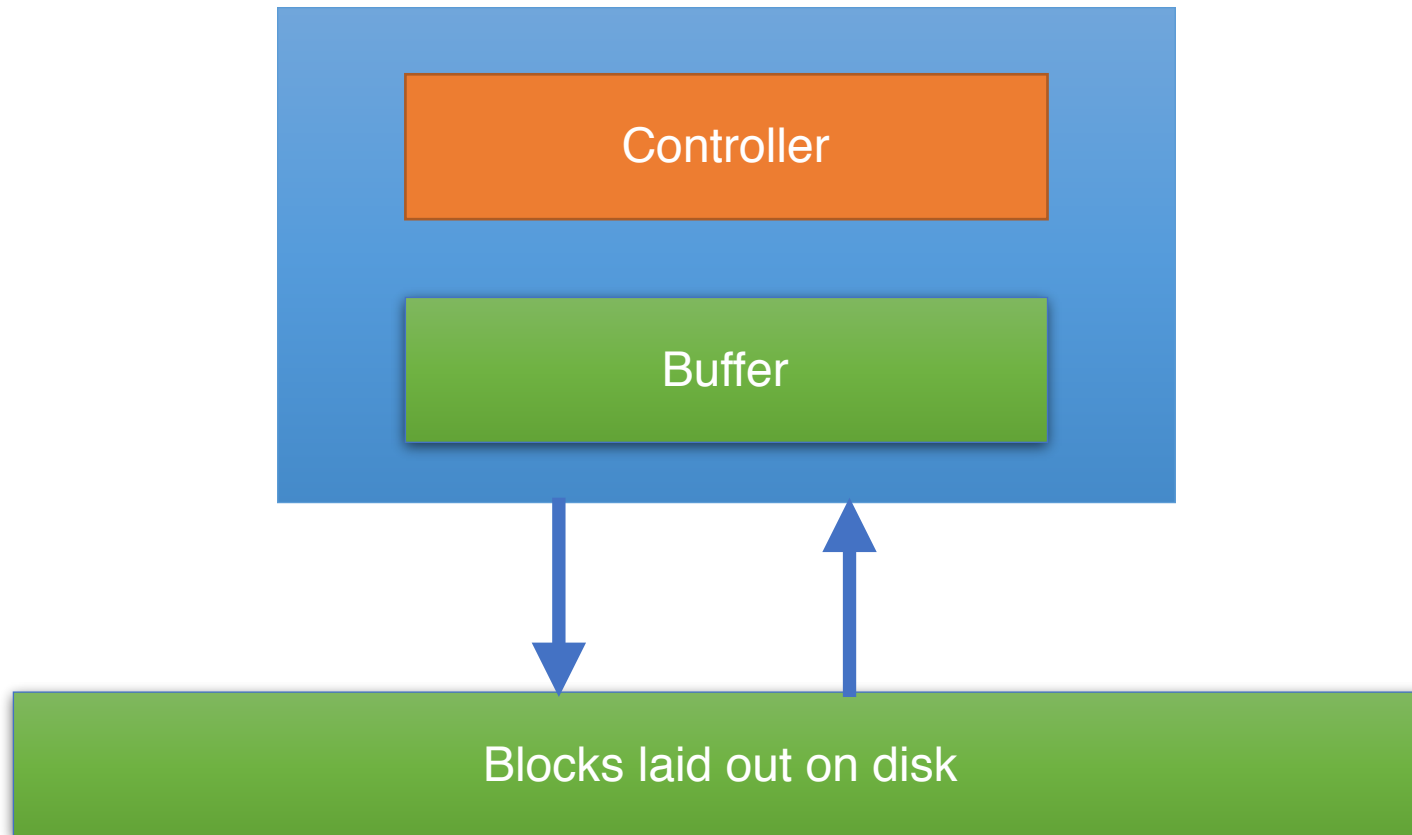
# Indexes Seem Magical...

- Q: what are downsides of indexes?
- A: the maintenance costs during updates, deletes, inserts
- If an index is not maintained, it may lead to inaccurate query results
- e.g., if a new tuple is inserted but it is not reflected in the index, it won't be returned as part of queries






# Aside: A Mental Model for a Database



- Data is laid out on blocks/ pages on stable storage (e.g., disk)
- Each block may have many 100s-1000s of tuples
- Data is retrieved a block at a time
- Stored in main-memory buffers
- Written back to stable storage



# Aside: A Mental Model for a Database



- Usually, computation is cheap relative to I/O
- So main cost in SQL queries is the cost of fetching the blocks and writing them out



# OK, Going Back to Indexes...

- Which indexes should we create?
- Answer: it depends!
- Many databases automatically create indexes for PRIMARY KEY or UNIQUE attributes. Several reasons:
  - It also helps enforcing the constraints! Updating the index can serve as a check for the constraint
  - Many queries lookup based on these attributes
    - `SELECT * FROM Film, FilmActor WHERE Film.film_id = FilmActor.film_id`
  - It is a high-cardinality attribute (as many values as number of tuples!)



# Why does cardinality matter?

- Databases retrieve data as “blocks”
- Say you are indexing based on attribute  $A$ , with 100 distinct values. You have 10,000 tuples. Each block contains 100 tuples.
- Q: How many blocks do we have?
- A: 100 blocks
- Assuming that the tuples with a given value of  $A$  ( $A = a$ ) are randomly distributed, every block, on average, every block has one tuple with  $A=a$ .
- Thus, we can't avoid looking at all blocks  $\Rightarrow$  the index doesn't help.
- Q: What if we had a 1000 distinct values instead of 100? How many blocks would we have to look at then?
- A: 10 blocks, so there may be some savings.



# Why does cardinality matter?

- Databases retrieve data as “blocks”
- Say you are indexing based on attribute A, with 100 distinct values. You have 10,000 tuples. Each block contains 100 tuples.
- However, if you know anything about disks, you will know that reading sequentially is much much faster than jumping across the disk
  - Even reading 1 in 10 blocks may not be enough to justify gains
  - Usually, 1 in 100 is enough to justify gains
- So, indexes on low cardinality attributes are not a great idea
- Exception: if data is “clustered” on that attribute
  - If we knew that all values of  $A = a$  are together (as opposed to randomly shuffled), we can simply read one block



# How Do We Decide?

- Good to create attributes used often in WHERE clauses
- Good to create on key attributes (but usually done by default)
- Good on attributes that are “clustered”
- Not very good on attributes where there are many tuples per value of attribute
- Not very good on tables that are modified more often than queried



# Types of Indexes: B+ Tree

- The most popular type of index in databases are B+Trees
- A generalization of binary trees
- Properties:
  - High fanout
    - Each “node” is a block
    - If “b” is fanout is  $\sim 1000$ ,  $b^4 = 10^{12}$
  - Self-balancing
    - Maintains a certain fill-factor (usually half) by splitting nodes and percolating changes to maintain balance
    - We won't worry about those algorithms in this class




At least halfway filled, except root

Find A = 30

Find A = [30,63]

Nodes



Pointers to records





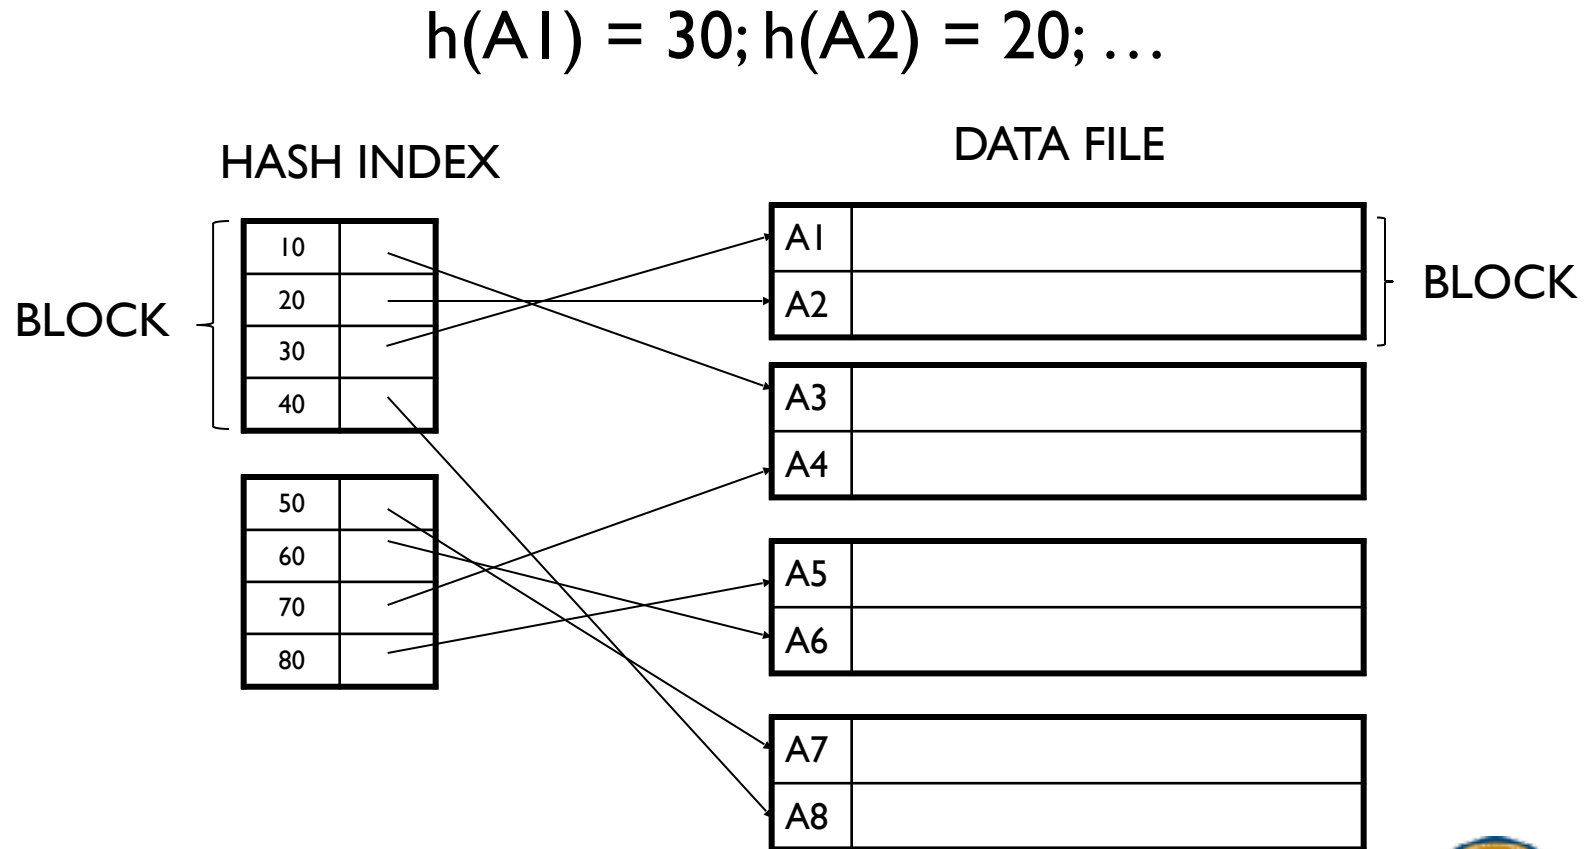
# Types of Indexes: B+ Tree

- Logarithmic complexity
  - For searching (as we saw)
  - For insert, delete, update (not covered)
- In practice, depth of tree = number of nodes (blocks) retrieved is small
- Can handle both value lookups as well as range lookups
- In PostgreSQL this is the default index type



# Types of Indexes: Hash Indexes

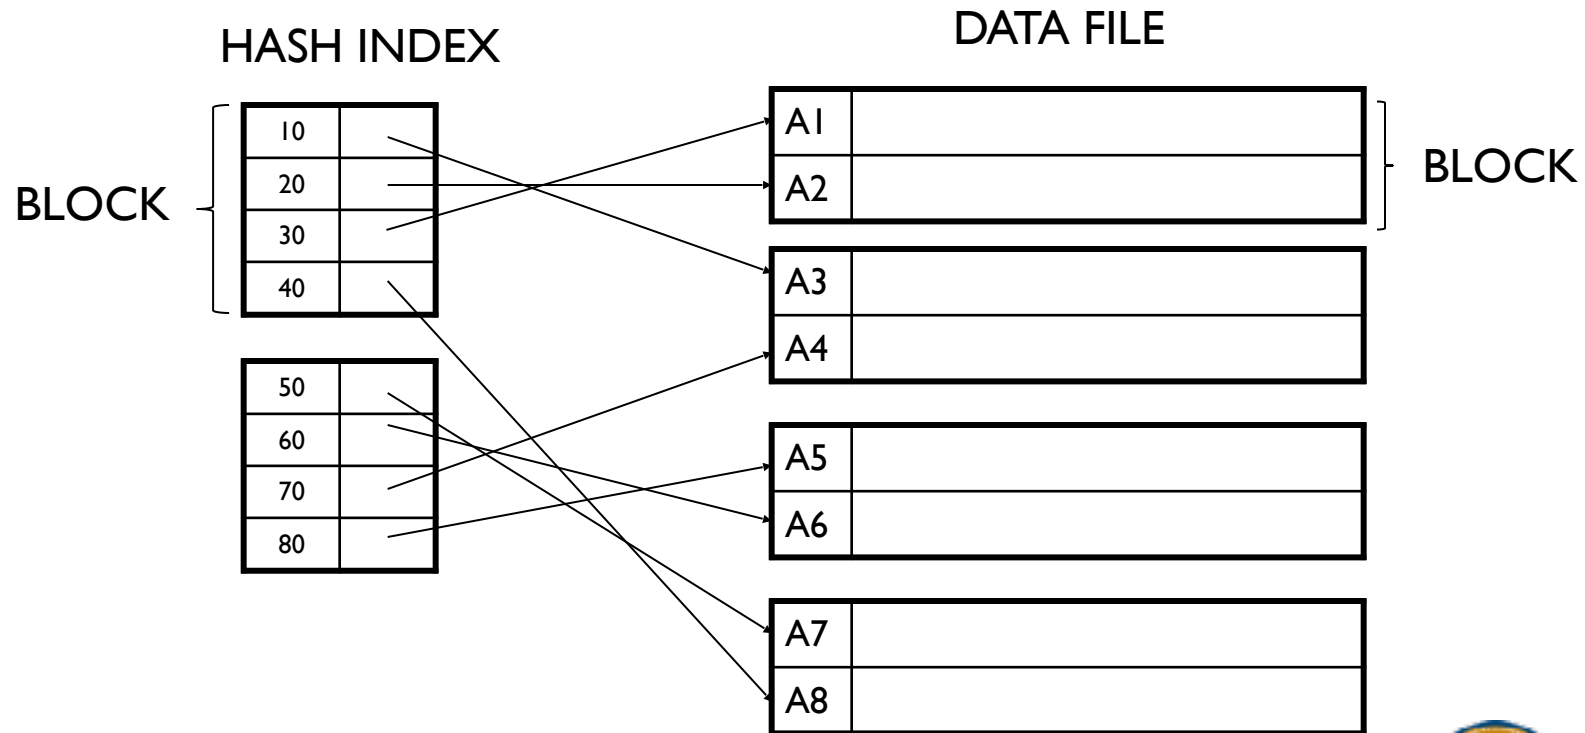
- Like hash-based dictionaries or maps in programming languages
  - A hash function applied to value  $a$ , i.e.,  $h(a)$  yields location where pointers for tuples with  $A = a$  are found
- Main challenge: how to “extend” the hash table as values of  $A$  grow
  - See linear or extendible hashing



# Types of Indexes: Hash Indexes

- Constant complexity  $O(1)$
- Can only handle value lookups, not range lookups
  - Why?
- In PostgreSQL done as follows:
  - **CREATE INDEX**  
**<name> ON <table>**  
**USING hash (column);**

$$h(A1) = 30; h(A2) = 20; \dots$$



# Other Types of Indexes

- Inverted Indexes: valuable for keyword search (more later, hopefully!)
  - In PostgreSQL, known as GIN indexes
- R-Tree: valuable for range-based lookups in k-dimensions
  - e.g., searching for things in rectangular range in a map
  - In PostgreSQL, a special case of a GIST index
- Lots of other types of indexes!
  - Partial indexes
  - Indexes on expressions
  - ...



# Indexes Summary

- Indexes are great!
- They speed up certain types of queries
- But we should use them carefully
  - Only on high cardinality = “selective” attributes like keys
  - Or on frequently queried but rarely updated attributes
- Lots of different types of indexes: but B+trees are most popular; hash indexes are a close second

