

Recap

- Data-savviness is the future!
- Notion of a DBMS
- The relational data model and algebra: bags and sets
- SQL Queries, Modifications, DDL
- Database Design
- Views, constraints, and triggers
- Indexes
- Query processing & optimization
- Next: transactions



So far...

- We have explored how to interact with a DBMS using SQL...
- However, at any time, many users can be interacting with the DBMS via various modalities...
- Imagine a bank... activities include:
 - Users may be making changes via mobile apps, websites, ATMs, ...
 - These changes may span multiple SQL statements (e.g., move money from one account to the next)
 - Banks may be updating investments, interests rates, ...
- Things could go wrong:
 - Users change their mind mid-sequence of changes
 - ATMs, mobile apps or even backend servers can crash



What are ways things can go wrong?

- **Example 1: interference**

- Flight Booking Program:

- A: Check if seat is available: SELECT-FROM-WHERE

- B: Book seat (change availability of seat from T to F, and assign passenger name to that seat): UPDATE-SET...

- Two simultaneous runs 1 & 2

- A1

- A2

- B1

- B2

- Two executions could lead to same seat being assigned to both passengers (but only user 2 has the seat actually)



What is the solution?

- Group the “SFV” (which verified availability) and “USV” (which reserved the seat) into one unit, called a TRANSACTION (or txns)
- *A transaction is a sequence of operations that are considered a basic unit within a database*
- Typical DBMSs guarantee *serializability* of txns
 - The state at any point is equivalent to a *serial* order of the txns
 - That is, either user 1’s txn executes first, and then user 2, or the other way around, rather than in parallel
 - In actuality, the DBMS doesn’t force the txns to execute in strict sequence, but uses intelligence to ensure correct parallelization
 - e.g., if the two users are reserving two different seats, it’s probably OK to parallelize



What else could go wrong?

- **Example 2: problems even with a single user**
- Even if one user is interacting with a DB at a time ...
- Transfer \$100 from checking to savings
 - Check if checking has $> \$100$
 - Reduce checking balance by \$100
 - Increase savings balance by \$100
- What if there's a crash after the second step?



What is the solution?

- Databases will ensure that the txn happens all at once or not at all
 - So either balance has been correctly transferred or not at all
- Q: How might a database do this?
- The way to maintain this is to “log” all changes that are being made so that the database can correctly *recover* from failures
 - For example, if we log that we’re going to make the change from checking to savings, we can consult the log to identify where we were when the program crashed
 - We can either “redo” the changes to persist them
 - Or “undo” the changes to reverse them



Declaring Txns

BEGIN;

SQL statement 1;

SQL statement 2;

...

SQL statement n;

COMMIT;

- SQL statements 1...n are an *atomic* unit.
- These atomic units are *serialized*.



Declaring Txns

BEGIN;

SQL statement l;

SAVEPOINT sl;

...

ROLLBACK TO sl;

SQL statement i;

...

SQL statement n;

COMMIT;

- ROLLBACK TO sl causes changes from sl onwards to be rolled back.
- A complete ROLLBACK will ensure that no effects from the txn are present



Txn Manager & Log Manager

- Components of the DBMS
- Job: making sure that txns execute as expected
- Purpose 1: log changes so that you can recover from failures
 - The *write-ahead log* allows you to write what you're planning to do before you do it so that you don't mess up the data, and you can recover from errors
- Purpose 2: make sure that simultaneous txns don't interfere with each other
 - The *concurrency control* module is used to ensure that txns are correctly parallelized
 - Various methods: locks, timestamps, ...
- Won't cover logging/recovery and concurrency control in this class...



Transaction Manager and Log Manager

- Together, achieve the so-called **ACID** properties
 - ATOMICITY: Each txn happens in entirety or not at all.
 - e.g., if a user's sequence of updates fails because their mobile app crashes, it either reflects entirely or not at all
 - *Via logging (and recovery)*
 - CONSISTENCY: Each txn leaves the database in a consistent state
 - e.g., if the database obeys constraints (ref integrity, NULL, PK etc.) prior to changes, they obey it after changes
 - *Via concurrency control (and checking of constraints prior to commit)*
 - ISOLATION: Each txn happens without interference from others
 - *Via concurrency control*
 - DURABILITY: Each txn's effects are durable, especially in the face of power or other types of failures
 - *Via logging (and recovery)*



Speeding Txns Up

- Maintaining these four properties over the course of many simultaneous txns is hard!
- Can cause delays ...
- As a user of txns, there are several things we may do to speed things up by indicating properties of a txn or relaxing requirements ...



I. Indicating Properties

- Can specify a txn as READ ONLY

SET TRANSACTION READ ONLY;

- Q: Why might it be helpful to indicate that a txn is read only?
- Benefit: DBMSs can use this to schedule this txn to happen in parallel with other READ ONLY txns, since there is no danger of conflicts (and therefore loss of isolation)
- Otherwise, default is READ WRITE



2. Relaxing Requirements

- Can relax requirements by setting an ISOLATION LEVEL for a txn.
- Default is SERIALIZABLE
 - Effect is equivalent to a serial schedule
- To motivate various ISOLATION LEVELS, let's see how we can relax requirements to read uncommitted data



2. Relaxing Requirements: Dirty Reads

- A *dirty read* is a read of *dirty data*: data written by a txn that hasn't committed yet.
- T1: Transfer \$100 from checking to savings
 - Check if checking has >\$100
 - Reduce checking balance by \$100
 - Increase savings balance by \$100
- Q: Say another T2 is OK with “dirty reads”. What types of T2s might have problems?
- Say T2 simply prints the sum of checking and savings and reads the balances after step 2
 - But this might be OK! Customer might just check again, and T1's effects might be completed.



2. Relaxing Requirements: Dirty Reads

- A *dirty read* is a read of *dirty data*: data written by a txn that hasn't committed yet.
- T1: Transfer \$100 from checking to savings
 - Check if checking has $> \$100$
 - Reduce checking balance by \$100
 - Increase savings balance by \$100
- Say T2 simply prints the sum of checking and savings
- Say T3 provides a bonus based on if checking + savings $> \$10000$
 - Customer may be less happy with an incorrect answer



2. Relaxing Requirements: Dirty Reads

- You can allow a transaction to read dirty data via the following:

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

- Note that an ISOLATION LEVEL is from the perspective of a txn:
 - That is, the given txn is reading dirty data
 - Not a global property of all txns



2. Relaxing Requirements: Read Committed

- The *read committed* isolation level ensures that a transaction only reads committed data.
- T1: Transfer \$100 from checking to savings
 - Check if checking has >\$100
 - Reduce checking balance by \$100
 - Increase savings balance by \$100
- Say T2 (as before) simply prints the sum of checking and savings
 - Q: If T2 is run as “read committed”, will it have any issues?



2. Relaxing Requirements: Read Committed

- The *read committed* isolation level ensures that a transaction only reads committed data.
- Done by

SET TRANSACTION READ WRITE
ISOLATION LEVEL READ COMMITTED;

- One issue is that over the course of a txn, it might see different instances of committed data for the same query (based on different txns that may have committed in the interim)



2. Relaxing Requirements: Read Committed

- One issue is that over the course of a txn, it might see different instances of committed data for the same query (based on different txns that may have committed in the interim)
- T1: Transfer \$100 from checking to savings
 - Check if checking has $> \$100$
 - Reduce checking balance by \$100
 - Increase savings balance by \$100
- Q: Can you come up with a T2 that will have issues with T1 if T2 does “read committed”
- A: T2 that computes the sum of checking and savings and splits up the reading of checking and savings across two steps could straddle T1 “before” and “after”



2. Relaxing Requirements: Repeatable Read

- Next isolation level: REPEATABLE READ

SET TRANSACTION READ WRITE

ISOLATION LEVEL REPEATABLE READ;

- Using the REPEATABLE READ isolation level, if the txn sees a certain tuple for a given query, it will continue to see the same tuple if query is issued again
 - So a tuple cannot “disappear” over the course of the transaction
 - One issue is that the query may also see new “phantom” data
- Sadly, REPEATABLE READ doesn't quite help with our example.



ISOLATION LEVELs

SET TRANSACTION READ WRITE

ISOLATION LEVEL READ UNCOMMITTED;

SET TRANSACTION READ WRITE

ISOLATION LEVEL READ COMMITTED;

SET TRANSACTION READ WRITE

ISOLATION LEVEL REPEATABLE READ;

SET TRANSACTION READ WRITE

ISOLATION LEVEL SERIALIZABLE; // Default



Txns and Recovery: Summary

- The database handles many users concurrently interacting with the data as well as failures and errors in a seamless way
- However, this handling can come with overheads, reducing # of txns serviced per second
- Can regulate this by various things:
 - Setting txns as read-only
 - Keeping them “small”
 - Setting more tolerant isolation levels
 - Don’t leave a txn “idle” longer than necessary
 - Use a connection pooling system (to reduce overhead)

