

Some Reminders for a Seamless Online Class...

- Please turn on your video
- Mute yourself (press and hold spacebar when you'd like to talk)
- Don't do anything you wouldn't do in an in-person class
- I will occasionally check the chat for messages if you'd like to share there instead
- Please say your name before you speak



Logistics

- Midterm and Programming Deadline Extended
 - Midterm report next
- Zoom can record! We'll share those videos (please don't share those broadly — just for private consumption)
- Any other questions?



Midterm Project Report (now due 3/30)

- From the FAQs:
 - “What goes in the midterm report?”

The midterm report is a 4-page summary outlining the project goals, its relevance to the course, key milestones/timeline, proposed design and architecture, the name of your project partner, and preliminary work done on the project. The midterm report will be evaluated based on the project description and proposed approach, justification of design and architecture, understanding of related work in this space, the feasibility of the projected plan, as well as promising results shown in the preliminary findings.



Recap

- Data-savviness is the future!
- “Classical” relational databases
 - Notion of a DBMS
 - The relational data model and algebra: bags and sets
 - SQL Queries, Modifications, DDL
 - Database Design
 - Views, constraints, triggers, and indexes
 - Query processing & optimization
 - Transactions
- Non-classical data systems
 - Semi-structured data and document stores
 - Next: unstructured data and search engines



Unstructured data

- Today, we will focus our attention on unstructured data that doesn't really have any “structure”
 - Unlike semistructured data, which does have some structure
 - It's a nested, flexible structure
- Unstructured data comes in many forms:
 - Text, images, video, audio
 - Our focus today is text — text is pretty general purpose:
 - Can capture “metadata” or “descriptions” for images and video
 - Can capture transcriptions for sound in video and audio too

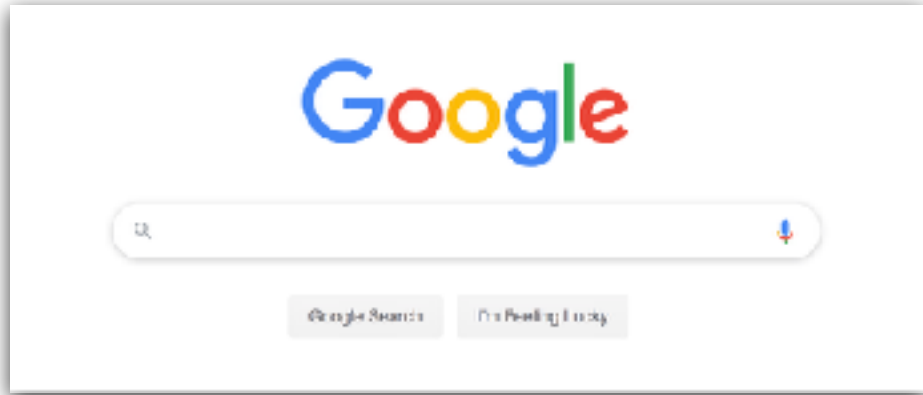


Querying Unstructured Data

- We've all queried unstructured data: search engines!
 - The field of designing better, more efficient search engines is called “information retrieval”
- We're going to get a bird's eye view of how to build a search engine
 - ...



First: the interface



- A simple text (or keyword) search field
- Sequences of keywords
- Can be more advanced:
 - Can specify exact word or phrase using “”
 - “Jack Ryan”
 - Can specify one word or the other using OR
 - Jack OR Ryan
 - Can specify a word or phrase not appearing
 - -“Jack Ryan”
 - Can specify file type, region, language, domain, ...




Not Just Restricted to Internet Search

- Product search
 - Amazon, Ebay, Walmart, ...
- Entertainment search
 - Netflix, Spotify, Youtube, ...
- Social media search
 - Twitter, Facebook, ...
- Email search



Interfaces Can Vary: Refining User Intent via Knobs




- Search facets are ways of organizing the results for a search query
 - Allowing users to browse the results using various “facets” or properties
- Search criteria can be used to reorder the search results



Interfaces Can Vary: Refining User Intent via Query Modification

- Search engines also suggest variations or completions of queries to assist users



Interfaces Can Vary: Structured Information

The screenshot shows a Google search for "jack ryan". The search results include a knowledge panel for "Tom Clancy's Jack Ryan" with information about its availability on streaming services (Amazon Prime Video, YouTube, Google Play, YouTube) and a Wikipedia entry for the TV series. A blue callout box points to the knowledge panel with the text "Structured data from knowledge bases".

Structured data from knowledge bases

- Knowledge bases (KBs) store information about various popular real-world entities
 - Can be used to provide structured information about these entities for quick lookup
 - Attributes like the fact that it is a movie, what its IMDB score, etc.
- Imagine a KB to be a hierarchy of entities
- Most search engines maintain KBs and use it to enhance their search results



Why the fuss?

- Q: Why do we have all these modules for understanding user intent better? Why not a simple text field and that's it?
 - After all, SQL is just “one query”, and we have no autocompletion, refinement, faceting, ...
- Considerable ambiguity in what the user wants and what they may find for a given text search query
 - e.g., Jack Ryan could refer to the books or the TV shows, but they may not get the books if they just typed “Jack Ryan”
 - Or the document/item could be an inexact match
- These are used by end-users, not by programmers or applications
 - If a web form is triggering a SQL query, this is usually rigorously tested



Formal Setting for Information Retrieval

- Fixed input: a set of documents
- Variable input: a search keyword query
- Goal:
 - Retrieve documents that “match” the keyword query (or underlying information need)
- Metrics
 - Precision: fraction of documents that “match” returned out of the number of documents retrieved
 - Recall: fraction of documents that “match” returned out of the total number of documents that “match”
 - Often, when there are lots of matches, these metrics become somewhat meaningless
 - “Precision@k”: proportion of retrieved documents that “match” out of first k



OK... How do we make this happen

- Let's talk about the data systems underlying search interfaces
- For now, simplifying the query:
 - “a” AND “b” AND NOT “c”
 - “Clinton” AND “Bush” AND NOT “Obama”
 - (I want to read articles that say compare Clinton to Bush, without talking about all three)
- Say we want to do this across a bunch of news articles (documents).
 - And we want to find all the matches
- One option: going through each document and checking what the word is — very slow!



Another Approach: Term-document incidence matrix

- Precomputed matrix
- Simply look at rows corresponding to Obama, Clinton, Bush.
- Perform binary AND of Clinton entry, Bush entry, and negation of Obama entry
- Q: why is this not ideal?
- A: most of this matrix would be empty!

	Doc 1	Doc 2	Doc 3	...
Obama	0	0	1	...
Clinton	1	1	1	...
Bush	0	1	1	...
...



Empty Term-Document Incidence Matrix

- If we have 1000 documents (each of length 1000 words) and there are 1M words, matrix has $1000 \times 1M = 1B$ entries
- Q: But how many 1s does the matrix have (in worst case)?
 - $1000 \times 1000 = 1M$
- Q: What can we do?
 - Compression!

	Doc 1	Doc 2	Doc 3	...
Obama	0	0	1	...
Clinton	1	1	1	...
Bush	0	1	1	...
...



Compressing the Matrix

- If we have 1000 documents (each of length 1000 words) and there are 1M words, matrix has $1000 \times 1M = 1B$ entries
- Compression approach:
 - For each word, maintain **sorted** list of documents that contain the word (more later) - called a **posting list**
 - Q: total size of posting list?
 - 1M entries, each of 10 bits
 - 10M

	Doc 1	Doc 2	Doc 3	...
Obama	0	0	1	...
Clinton	1	1	1	...
Bush	0	1	1	...
...

How do Posting Lists Help?

- Obama -> 10, 24, 125, 259, 1025, 2314, ...
- Bush -> 10, 15, 17, 259, 2001, 2547, ...
- Clinton -> 10, 15, 24, 17, 125, 1005, 2001, 2347, ...
- What approach would we use to find “Clinton” AND “Bush” AND NOT “Obama”?
- Essentially a “Merge”!



Overall Structure

- A collection of postings lists, one per word is also called an **inverted index**
- Thus, an inverted index allows us to quickly retrieve all matching documents for queries involving boolean connectives between words
- Q:What are ways this is insufficient?



Improving the Basic Process

- Inexact word specification
 - “modify” vs. “modified” vs. “modification” (roughly same meaning)
 - Stemming: replace all of “modify” “modifying” “modified” with “modif”
 - IL and Illinois, USA and United States
 - Canonicalization: pick a canonical form
- Some words are very frequent “to”, “and” etc.
 - These are called **stop words**
 - These words are omitted from the matching
 - Their posting lists will be very long, and a lot of time will be wasted on them!



Improving the Basic Process

- Wanting the words to be close to each other
 - “jack” AND “ryan” may match an article about “jack black” and “ryan adams” (but you want them to be near each other!)
 - Q: how may we able to capture this?
 - A: post-processing may be really bad
 - Embed information about positions into posting list
 - Jack -> 10 (1, 10, 20), 24 (1, 15, 27), ...
 - Ryan -> 10 (2, 15, 30), 24 (3, 25, 57), ...
 - Or do “biword” posting lists
 - Can extend to many words A1 A2 A3 A4
 - A1 A2 and A2 A3 and A3 A4



Scoring the Matches

- Say you found 100s of documents that matched the query.
- Now we need to rank it to show it to the user
- Typically use some combination of:
 - The inherent “quality” of a document
 - If on the web, this could be based on webpages that link to the webpage — page rank
 - The importance of the term to the document
 - One way to score this is called TF-IDF
 - Term Frequency X Inverse Document Frequency

$$\text{\# of term occ. in doc} \quad \frac{f_t}{\sum_{t'} f_{t'}} \times \log \frac{N}{N_t} \quad \text{\# of docs with term}$$



Generating the Inverted Index

- Phase 1:
 - For each document, generate a “canonicalized term, documentID” pair for all non-stop-word terms in the document
 - Can encode more information with the pair, e.g., list of locations, frequency of occurrence, etc.
- Phase 2:
 - Sort these pairs based on the term, documentID pair
 - May need to use two or multi-pass sort (remember query processing!)
- Phase 3:
 - Read in the sorted pairs based on term, concatenate documentID together to form a posting list for that term
 - (We may be able to merge Phase 3 with Phase 2)



Examples of Data Systems

- Lucene
 - Open source text indexing and search library (since 2000s)
- ElasticSearch
 - Combines Lucene-like text indexing/search with JSON querying capabilities (since 2010)
- Database systems also support text indexing and search
 - For example, PostgreSQL supports
 - GIN indexes (Generalized INverted Indexes) for querying text
 - Basic string stemming, canonicalization, tokenization capabilities



Summary

- Unstructured data is everywhere
- The discipline of retrieving relevant unstructured data is called information retrieval
- Keyword search is the primary way to query unstructured data
 - Terms with boolean connectives
- Many variations in terms of the keyword search interface itself
- Inverted indexes as a key mechanism to support keyword search
- When coupled with various scoring mechanisms provide a good foundations for searching unstructured data

